

MIS 583 Assignment 5: YOLO Object Detection on PASCAL VOC

Before we start, please put your name and SID in following format:

: LASTNAME Firstname, ?00000000 // e.g.) 李晨愷 M114020035

Your Answer:

Hi I'm 游雅淇, B104020012.

Google Colab Setup

Next we need to run a few commands to set up our environment on Google Colab. If you are running this notebook on a local machine you can skip this section.

Run the following cell to mount your Google Drive. Follow the link, sign in to your Google account (the same account you used to store this notebook!) and copy the authorization code into the text box that appears below.

```
In [ ]: from google.colab import drive
drive.mount('/content/drive', force_remount=True)
```

Mounted at /content/drive

How to Get Data

請先到共用雲端硬碟將檔案 `VOCdevkit_2007.zip`，建立捷徑到自己的雲端硬碟中。

操作步驟

1. 點開雲端[連結](#)
2. 點選右上角「新增雲端硬碟捷徑」
3. 點選「我的雲端硬碟」
4. 點選「新增捷徑」

完成以上流程會在你的雲端硬碟中建立一個檔案的捷徑，接著我們在colab中取得權限即可使用。

Unzip Data

解壓縮 `VOCdevkit_2007.zip`

- `VOC2007` : 包含了train/val的所有圖片
- `VOC2007test` : 包含了test的所有圖片

其中 `train` 的圖片 3756 張，`val` 的圖片 1255 張，`test` 的圖片 4950 張。

注意: 若有另外設定存放在雲端硬碟中的路徑，請記得本處路徑也須做更動。

Notice: Please put "VOCdevkit_2007" folder under data folder.

```
In [ ]: !unzip -qq ./drive/MyDrive/DeepLearning/A5/data/VOCdevkit_2007.zip
```

```
In [ ]: %cd ./drive/MyDrive/A5
```

```
[Errno 2] No such file or directory: './drive/MyDrive/A5'
/content/drive/MyDrive/A5
```

Import package

```
In [ ]: import os
import random

import cv2
import numpy as np

import csv

import torch
from torch.utils.data import DataLoader
from torchvision import models

from src.resnet_yolo import resnet50
from src.densenet_yolo import densenet121
from yolo_loss import YoloLoss
from src.dataset import VocDetectorDataset
from src.eval_voc import evaluate, test_evaluate
from src.predict import predict_image
from src.config import VOC_CLASSES, COLORS
from kaggle_submission import write_csv

import matplotlib.pyplot as plt
import collections

%matplotlib inline
%load_ext autoreload
%autoreload 2
```

The autoreload extension is already loaded. To reload it, use:
%reload_ext autoreload

Initialization

```
In [ ]: device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
```

```
In [ ]: # YOLO network hyperparameters
B = 2 # number of bounding box predictions per cell
S = 14 # width/height of network output grid (larger than 7x7 from paper s
```

To implement Yolo we will rely on a pretrained classifier as the backbone for our detection network. PyTorch offers a variety of models which are pretrained on ImageNet in the `torchvision.models` package. In particular, we will use the ResNet50 architecture as a base for our detector. This is different from the base architecture in the Yolo paper and also results in a different output grid size (14x14 instead of 7x7).

Models are typically pretrained on ImageNet since the dataset is very large (> 1 million images) and widely used. The pretrained model provides a very useful weight initialization for our detector, so that the network is able to learn quickly and effectively.

```
In [ ]: load_network_path = None #'checkpoints/best_detector.pth'
pretrained = True

# use to load a previously trained network
if load_network_path is not None:
    print('Loading saved network from {}'.format(load_network_path))
    net = densenet121(S=S).to(device)
    net.load_state_dict(torch.load(load_network_path))
else:
    print('Load pre-trained model')
    net = densenet121(pretrained=pretrained, S=S).to(device)
```

Load pre-trained model

```
/usr/local/lib/python3.10/dist-packages/torchvision/models/_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the future, please use 'weights' instead.
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed in the future. The current behavior is equivalent to passing `weights=DenseNet121_Weights.IMAGENET1K_V1`. You can also use `weights=DenseNet121_Weights.DEFAULT` to get the most up-to-date weights.
  warnings.warn(msg)
```

```
In [ ]: learning_rate = 0.001
num_epochs = 60
batch_size = 24

# Yolo loss component coefficients (as given in Yolo v1 paper)
lambda_coord = 5
lambda_noobj = 0.5
```

Reading Pascal Data

The train dataset loader also using a variety of data augmentation techniques including random shift, scaling, crop, and flips. Data augmentation is slightly more complicated for detection datasets since the bounding box annotations must be kept consistent throughout the transformations.

Since the output of the detector network we train is an $S \times S \times (B \times 5 + C)$, we use an encoder to convert the original bounding box coordinates into relative grid bounding box coordinates corresponding to the expected output. We also use a decoder which allows us to convert the opposite direction into image coordinate bounding boxes.

Notice: Please put "VOCdevkit_2007" folder under data folder.

```
In [ ]: file_root_train = 'data/VOCdevkit_2007/VOC2007/JPEGImages/'
annotation_file_train = 'data/voc2007train.txt'

train_dataset = VocDetectorDataset(root_img_dir=file_root_train, dataset_file=
train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True,
print('Loaded %d train images' % len(train_dataset))
```

Initializing dataset
Loaded 3756 train images

```
In [ ]: file_root_val = 'data/VOCdevkit_2007/VOC2007/JPEGImages/'
        annotation_file_val = 'data/voc2007val.txt'

        val_dataset = VocDetectorDataset(root_img_dir=file_root_val, dataset_file=an
        val_loader = DataLoader(val_dataset, batch_size=batch_size, shuffle=False, num
        print('Loaded %d val images' % len(val_dataset))
```

Initializing dataset
Loaded 1255 val images

```
In [ ]: data = train_dataset[0]
```

Set up training tools

```
In [ ]: criterion = YoloLoss(S, B, lambda_coord, lambda_noobj)
        optimizer = torch.optim.SGD(net.parameters(), lr=learning_rate, momentum=0.9)
```

```
In [ ]: # data normalization
```

Train detector

```
In [ ]: best_val_loss = np.inf
        learning_rate = 1e-3
        for epoch in range(num_epochs):
            net.train()

            # Update learning rate late in training
            if epoch == 30 or epoch == 40:
                learning_rate /= 10.0

            for param_group in optimizer.param_groups:
                param_group['lr'] = learning_rate

            print('\n\nStarting epoch %d / %d' % (epoch + 1, num_epochs))
            print('Learning Rate for this epoch: {}'.format(learning_rate))

            total_loss = collections.defaultdict(int)

            for i, data in enumerate(train_loader):
                data = (item.to(device) for item in data)
                images, target_boxes, target_cls, has_object_map = data
                pred = net(images)
                loss_dict = criterion(pred, target_boxes, target_cls, has_object_map)
                for key in loss_dict:
                    total_loss[key] += loss_dict[key].item()

                optimizer.zero_grad()
                loss_dict['total_loss'].backward()
                optimizer.step()

                if (i+1) % 50 == 0:
                    outstring = 'Epoch [%d/%d], Iter [%d/%d], Loss: ' % ((epoch+1, num_epochs), i+1, i+1)
                    outstring += ', '.join( "%s=%.3f" % (key[:5], val / (i+1)) for key, val in loss_dict.items())
                    print(outstring)

            # evaluate the network on the val data
            if (epoch + 1) % 5 == 0:
```

```
val_aps = evaluate(net, val_dataset_file=annotation_file_val, img_ro
print(epoch, val_aps)
with torch.no_grad():
    val_loss = 0.0
    net.eval()
    for i, data in enumerate(val_loader):
        data = (item.to(device) for item in data)
        images, target_boxes, target_cls, has_object_map = data

        pred = net(images)
        loss_dict = criterion(pred, target_boxes, target_cls, has_object
        val_loss += loss_dict['total_loss'].item()
    val_loss /= len(val_loader)

if best_val_loss > val_loss:
    best_val_loss = val_loss
    print('Updating best val loss: %.5f' % best_val_loss)
    torch.save(net.state_dict(), 'checkpoints/best_detector.pth')

if (epoch+1) in [5, 10, 20, 30, 40, 50, 60, 70, 80]:
    torch.save(net.state_dict(), 'checkpoints/detector_epoch_%d.pth' % (e

torch.save(net.state_dict(), 'checkpoints/detector.pth')
```

Starting epoch 1 / 60
Learning Rate for this epoch: 0.001
Epoch [1/60], Iter [50/157], Loss: total=23.284, reg=0.652, containing_obj=1.509, no_obj=25.329, cls=5.851
Epoch [1/60], Iter [100/157], Loss: total=15.448, reg=0.566, containing_obj=1.711, no_obj=13.289, cls=4.264
Epoch [1/60], Iter [150/157], Loss: total=12.426, reg=0.515, containing_obj=1.730, no_obj=9.159, cls=3.540
Updating best val loss: 6.67858

Starting epoch 2 / 60
Learning Rate for this epoch: 0.001
Epoch [2/60], Iter [50/157], Loss: total=6.055, reg=0.423, containing_obj=1.637, no_obj=0.978, cls=1.815
Epoch [2/60], Iter [100/157], Loss: total=5.860, reg=0.416, containing_obj=1.577, no_obj=0.981, cls=1.710
Epoch [2/60], Iter [150/157], Loss: total=5.648, reg=0.406, containing_obj=1.521, no_obj=0.978, cls=1.610
Updating best val loss: 5.51628

Starting epoch 3 / 60
Learning Rate for this epoch: 0.001
Epoch [3/60], Iter [50/157], Loss: total=4.940, reg=0.372, containing_obj=1.364, no_obj=0.963, cls=1.235
Epoch [3/60], Iter [100/157], Loss: total=4.970, reg=0.376, containing_obj=1.363, no_obj=0.946, cls=1.253
Epoch [3/60], Iter [150/157], Loss: total=4.935, reg=0.378, containing_obj=1.349, no_obj=0.936, cls=1.226
Updating best val loss: 5.00272

Starting epoch 4 / 60
Learning Rate for this epoch: 0.001
Epoch [4/60], Iter [50/157], Loss: total=4.414, reg=0.347, containing_obj=1.258, no_obj=0.850, cls=0.994
Epoch [4/60], Iter [100/157], Loss: total=4.452, reg=0.354, containing_obj=1.275, no_obj=0.845, cls=0.986
Epoch [4/60], Iter [150/157], Loss: total=4.423, reg=0.350, containing_obj=1.268, no_obj=0.855, cls=0.979
Updating best val loss: 4.72996

Starting epoch 5 / 60
Learning Rate for this epoch: 0.001
Epoch [5/60], Iter [50/157], Loss: total=4.063, reg=0.328, containing_obj=1.173, no_obj=0.856, cls=0.823
Epoch [5/60], Iter [100/157], Loss: total=4.123, reg=0.333, containing_obj=1.191, no_obj=0.840, cls=0.847
Epoch [5/60], Iter [150/157], Loss: total=4.165, reg=0.340, containing_obj=1.201, no_obj=0.835, cls=0.846
---Evaluate model on test samples---

100%|██████████| 1255/1255 [00:58<00:00, 21.48it/s]

```
---class aeroplane ap 0.5128774410218597---
---class bicycle ap 0.4080926689793873---
---class bird ap 0.3591520464190196---
---class boat ap 0.14623523566161445---
---class bottle ap 0.2023822459299064---
---class bus ap 0.19506493506493505---
---class car ap 0.4892260661259066---
---class cat ap 0.4284075116376122---
---class chair ap 0.19844370539944864---
---class cow ap 0.12310142623441855---
---class diningtable ap 0.0--- (no predictions for this class)
---class dog ap 0.46758777962266385---
---class horse ap 0.38670109531706004---
---class motorbike ap 0.4152654665198618---
---class person ap 0.3772111220540883---
---class pottedplant ap 0.09845305168941854---
---class sheep ap 0.12333593474557089---
---class sofa ap 0.31616105358795416---
---class train ap 0.5467133637278708---
---class tvmonitor ap 0.5095881915586921---
---map 0.31520001706486445---
4 [0.5128774410218597, 0.4080926689793873, 0.3591520464190196, 0.1462352356
6161445, 0.2023822459299064, 0.19506493506493505, 0.4892260661259066, 0.428
4075116376122, 0.19844370539944864, 0.12310142623441855, 0.0, 0.46758777962
266385, 0.38670109531706004, 0.4152654665198618, 0.3772111220540883, 0.0984
5305168941854, 0.12333593474557089, 0.31616105358795416, 0.546713363727870
8, 0.5095881915586921]
Updating best val loss: 4.53227
```

Starting epoch 6 / 60

Learning Rate for this epoch: 0.001

Epoch [6/60], Iter [50/157], Loss: total=3.913, reg=0.316, containing_obj=1.162, no_obj=0.833, cls=0.757

Epoch [6/60], Iter [100/157], Loss: total=3.993, reg=0.326, containing_obj=1.173, no_obj=0.846, cls=0.768

Epoch [6/60], Iter [150/157], Loss: total=3.938, reg=0.323, containing_obj=1.153, no_obj=0.845, cls=0.748

Updating best val loss: 4.33410

Starting epoch 7 / 60

Learning Rate for this epoch: 0.001

Epoch [7/60], Iter [50/157], Loss: total=3.736, reg=0.306, containing_obj=1.167, no_obj=0.782, cls=0.648

Epoch [7/60], Iter [100/157], Loss: total=3.766, reg=0.309, containing_obj=1.157, no_obj=0.802, cls=0.662

Epoch [7/60], Iter [150/157], Loss: total=3.795, reg=0.311, containing_obj=1.154, no_obj=0.804, cls=0.686

Updating best val loss: 4.31756

Starting epoch 8 / 60

Learning Rate for this epoch: 0.001

Epoch [8/60], Iter [50/157], Loss: total=3.353, reg=0.274, containing_obj=1.045, no_obj=0.765, cls=0.558

Epoch [8/60], Iter [100/157], Loss: total=3.574, reg=0.294, containing_obj=1.108, no_obj=0.785, cls=0.603

Epoch [8/60], Iter [150/157], Loss: total=3.625, reg=0.299, containing_obj=1.117, no_obj=0.795, cls=0.617

Updating best val loss: 4.16529

Starting epoch 9 / 60

```
Learning Rate for this epoch: 0.001
Epoch [9/60], Iter [50/157], Loss: total=3.354, reg=0.281, containing_obj=
1.029, no_obj=0.786, cls=0.529
Epoch [9/60], Iter [100/157], Loss: total=3.379, reg=0.282, containing_obj=
1.044, no_obj=0.781, cls=0.533
Epoch [9/60], Iter [150/157], Loss: total=3.435, reg=0.288, containing_obj=
1.046, no_obj=0.789, cls=0.555
```

Starting epoch 10 / 60

```
Learning Rate for this epoch: 0.001
Epoch [10/60], Iter [50/157], Loss: total=3.435, reg=0.292, containing_obj=
1.036, no_obj=0.840, cls=0.521
Epoch [10/60], Iter [100/157], Loss: total=3.411, reg=0.287, containing_obj
=1.045, no_obj=0.814, cls=0.526
Epoch [10/60], Iter [150/157], Loss: total=3.401, reg=0.286, containing_obj
=1.045, no_obj=0.798, cls=0.525
---Evaluate model on test samples---
```

```
100%|██████████| 1255/1255 [00:53<00:00, 23.55it/s]
```



```
---class aeroplane ap 0.5117104167634492---
---class bicycle ap 0.4980647210609129---
---class bird ap 0.5042727868691758---
---class boat ap 0.2892333445962617---
---class bottle ap 0.17754853700339224---
---class bus ap 0.41616254246683065---
---class car ap 0.6012487227559637---
---class cat ap 0.5790325932580195---
---class chair ap 0.2709576660209653---
---class cow ap 0.31085989601911446---
---class diningtable ap 0.1571410786536837---
---class dog ap 0.45423074727987073---
---class horse ap 0.5916310443076743---
---class motorbike ap 0.49200443073480005---
---class person ap 0.4719044011513074---
---class pottedplant ap 0.19203491605232606---
---class sheep ap 0.35983788486766405---
---class sofa ap 0.36259328127591345---
---class train ap 0.667016656307068---
---class tvmonitor ap 0.5282259543334494---
---map 0.4217855810888921---
9 [0.5117104167634492, 0.4980647210609129, 0.5042727868691758, 0.2892333445
962617, 0.17754853700339224, 0.41616254246683065, 0.6012487227559637, 0.579
0325932580195, 0.2709576660209653, 0.31085989601911446, 0.1571410786536837,
0.45423074727987073, 0.5916310443076743, 0.49200443073480005, 0.47190440115
13074, 0.19203491605232606, 0.35983788486766405, 0.36259328127591345, 0.667
016656307068, 0.5282259543334494]
Updating best val loss: 4.15944
```

```
Starting epoch 11 / 60
Learning Rate for this epoch: 0.001
Epoch [11/60], Iter [50/157], Loss: total=3.357, reg=0.286, containing_obj=
1.023, no_obj=0.814, cls=0.495
Epoch [11/60], Iter [100/157], Loss: total=3.324, reg=0.282, containing_obj
=1.019, no_obj=0.799, cls=0.496
Epoch [11/60], Iter [150/157], Loss: total=3.321, reg=0.282, containing_obj
=1.018, no_obj=0.794, cls=0.497
Updating best val loss: 4.09678
```

```
Starting epoch 12 / 60
Learning Rate for this epoch: 0.001
Epoch [12/60], Iter [50/157], Loss: total=3.191, reg=0.271, containing_obj=
0.997, no_obj=0.782, cls=0.448
Epoch [12/60], Iter [100/157], Loss: total=3.182, reg=0.269, containing_obj
=0.998, no_obj=0.761, cls=0.457
Epoch [12/60], Iter [150/157], Loss: total=3.220, reg=0.271, containing_obj
=1.004, no_obj=0.774, cls=0.473
Updating best val loss: 3.99937
```

```
Starting epoch 13 / 60
Learning Rate for this epoch: 0.001
Epoch [13/60], Iter [50/157], Loss: total=3.060, reg=0.258, containing_obj=
0.964, no_obj=0.770, cls=0.420
Epoch [13/60], Iter [100/157], Loss: total=3.061, reg=0.258, containing_obj
=0.958, no_obj=0.788, cls=0.417
Epoch [13/60], Iter [150/157], Loss: total=3.070, reg=0.259, containing_obj
=0.972, no_obj=0.776, cls=0.417
```

```
Starting epoch 14 / 60
Learning Rate for this epoch: 0.001
```

```
Epoch [14/60], Iter [50/157], Loss: total=3.089, reg=0.268, containing_obj=
0.957, no_obj=0.784, cls=0.399
Epoch [14/60], Iter [100/157], Loss: total=3.049, reg=0.263, containing_obj
=0.957, no_obj=0.766, cls=0.395
Epoch [14/60], Iter [150/157], Loss: total=3.023, reg=0.257, containing_obj
=0.957, no_obj=0.765, cls=0.396
Updating best val loss: 3.88821
```

Starting epoch 15 / 60

Learning Rate for this epoch: 0.001

```
Epoch [15/60], Iter [50/157], Loss: total=2.877, reg=0.249, containing_obj=
0.887, no_obj=0.782, cls=0.356
```

```
Epoch [15/60], Iter [100/157], Loss: total=2.927, reg=0.254, containing_obj
=0.907, no_obj=0.785, cls=0.358
```

```
Epoch [15/60], Iter [150/157], Loss: total=2.943, reg=0.254, containing_obj
=0.910, no_obj=0.781, cls=0.371
```

---Evaluate model on test samples---

```
100%|██████████| 1255/1255 [00:54<00:00, 23.10it/s]
```

```
---class aeroplane ap 0.6284671251340119---
---class bicycle ap 0.5969974880742045---
---class bird ap 0.4228659841927902---
---class boat ap 0.2487633505400162---
---class bottle ap 0.2733061446403895---
---class bus ap 0.49335326538103735---
---class car ap 0.641302560205872---
---class cat ap 0.7059173770855257---
---class chair ap 0.30890833363036674---
---class cow ap 0.48072676890874616---
---class diningtable ap 0.2195648073696854---
---class dog ap 0.5402492654819987---
---class horse ap 0.5187337783692301---
---class motorbike ap 0.4764289566992827---
---class person ap 0.5185821625434274---
---class pottedplant ap 0.19795422542672084---
---class sheep ap 0.2823588897208816---
---class sofa ap 0.44679944934241805---
---class train ap 0.708707190020227---
---class tvmonitor ap 0.5750144091675318---
---map 0.46425007659671824---
14 [0.6284671251340119, 0.5969974880742045, 0.4228659841927902, 0.248763350
5400162, 0.2733061446403895, 0.49335326538103735, 0.641302560205872, 0.7059
173770855257, 0.30890833363036674, 0.48072676890874616, 0.2195648073696854,
0.5402492654819987, 0.5187337783692301, 0.4764289566992827, 0.5185821625434
274, 0.19795422542672084, 0.2823588897208816, 0.44679944934241805, 0.708707
190020227, 0.5750144091675318]
Updating best val loss: 3.84921
```

Starting epoch 16 / 60

Learning Rate for this epoch: 0.001

Epoch [16/60], Iter [50/157], Loss: total=2.864, reg=0.244, containing_obj=0.878, no_obj=0.752, cls=0.390

Epoch [16/60], Iter [100/157], Loss: total=2.905, reg=0.247, containing_obj=0.908, no_obj=0.756, cls=0.382

Epoch [16/60], Iter [150/157], Loss: total=2.918, reg=0.250, containing_obj=0.912, no_obj=0.760, cls=0.373

Starting epoch 17 / 60

Learning Rate for this epoch: 0.001

Epoch [17/60], Iter [50/157], Loss: total=2.864, reg=0.249, containing_obj=0.891, no_obj=0.775, cls=0.341

Epoch [17/60], Iter [100/157], Loss: total=2.876, reg=0.251, containing_obj=0.882, no_obj=0.788, cls=0.347

Epoch [17/60], Iter [150/157], Loss: total=2.837, reg=0.246, containing_obj=0.877, no_obj=0.776, cls=0.342

Starting epoch 18 / 60

Learning Rate for this epoch: 0.001

Epoch [18/60], Iter [50/157], Loss: total=2.814, reg=0.244, containing_obj=0.890, no_obj=0.744, cls=0.333

Epoch [18/60], Iter [100/157], Loss: total=2.826, reg=0.244, containing_obj=0.898, no_obj=0.753, cls=0.331

Epoch [18/60], Iter [150/157], Loss: total=2.808, reg=0.240, containing_obj=0.887, no_obj=0.755, cls=0.342

Updating best val loss: 3.83189

Starting epoch 19 / 60

Learning Rate for this epoch: 0.001

Epoch [19/60], Iter [50/157], Loss: total=2.602, reg=0.220, containing_obj=

```
0.821, no_obj=0.759, cls=0.300  
Epoch [19/60], Iter [100/157], Loss: total=2.676, reg=0.227, containing_obj=  
=0.854, no_obj=0.757, cls=0.307  
Epoch [19/60], Iter [150/157], Loss: total=2.701, reg=0.230, containing_obj=  
=0.855, no_obj=0.756, cls=0.316
```

```
Starting epoch 20 / 60  
Learning Rate for this epoch: 0.001  
Epoch [20/60], Iter [50/157], Loss: total=2.640, reg=0.231, containing_obj=  
0.836, no_obj=0.728, cls=0.285  
Epoch [20/60], Iter [100/157], Loss: total=2.664, reg=0.233, containing_obj=  
=0.837, no_obj=0.741, cls=0.292  
Epoch [20/60], Iter [150/157], Loss: total=2.678, reg=0.233, containing_obj=  
=0.846, no_obj=0.743, cls=0.298  
---Evaluate model on test samples---
```

```
100%|██████████| 1255/1255 [00:54<00:00, 22.86it/s]
```

```
---class aeroplane ap 0.5567289778568962---
---class bicycle ap 0.5965211766735979---
---class bird ap 0.49615716511938074---
---class boat ap 0.28964822427301506---
---class bottle ap 0.25131081805156824---
---class bus ap 0.5165638597249916---
---class car ap 0.6456949209177731---
---class cat ap 0.681195657278251---
---class chair ap 0.3164439773856242---
---class cow ap 0.5048663153659372---
---class diningtable ap 0.33361115812189535---
---class dog ap 0.5726212548596231---
---class horse ap 0.6330027775791869---
---class motorbike ap 0.5746623093659794---
---class person ap 0.4784480102910313---
---class pottedplant ap 0.18515028107540848---
---class sheep ap 0.37224171059664546---
---class sofa ap 0.3834791306561304---
---class train ap 0.6594695062021972---
---class tvmonitor ap 0.6712240022787418---
---map 0.48595206168369376---
19 [0.5567289778568962, 0.5965211766735979, 0.49615716511938074, 0.28964822
427301506, 0.25131081805156824, 0.5165638597249916, 0.6456949209177731, 0.6
81195657278251, 0.3164439773856242, 0.5048663153659372, 0.3336111581218953
5, 0.5726212548596231, 0.6330027775791869, 0.5746623093659794, 0.4784480102
910313, 0.18515028107540848, 0.37224171059664546, 0.3834791306561304, 0.659
4695062021972, 0.6712240022787418]
```

Starting epoch 21 / 60

Learning Rate for this epoch: 0.001

Epoch [21/60], Iter [50/157], Loss: total=2.593, reg=0.222, containing_obj=0.811, no_obj=0.734, cls=0.306

Epoch [21/60], Iter [100/157], Loss: total=2.633, reg=0.229, containing_obj=0.819, no_obj=0.745, cls=0.295

Epoch [21/60], Iter [150/157], Loss: total=2.628, reg=0.227, containing_obj=0.826, no_obj=0.746, cls=0.294

Starting epoch 22 / 60

Learning Rate for this epoch: 0.001

Epoch [22/60], Iter [50/157], Loss: total=2.672, reg=0.226, containing_obj=0.855, no_obj=0.787, cls=0.291

Epoch [22/60], Iter [100/157], Loss: total=2.576, reg=0.218, containing_obj=0.823, no_obj=0.756, cls=0.283

Epoch [22/60], Iter [150/157], Loss: total=2.561, reg=0.219, containing_obj=0.812, no_obj=0.745, cls=0.279

Updating best val loss: 3.80138

Starting epoch 23 / 60

Learning Rate for this epoch: 0.001

Epoch [23/60], Iter [50/157], Loss: total=2.501, reg=0.212, containing_obj=0.802, no_obj=0.719, cls=0.277

Epoch [23/60], Iter [100/157], Loss: total=2.467, reg=0.212, containing_obj=0.794, no_obj=0.707, cls=0.262

Epoch [23/60], Iter [150/157], Loss: total=2.525, reg=0.215, containing_obj=0.814, no_obj=0.723, cls=0.273

Starting epoch 24 / 60

Learning Rate for this epoch: 0.001

Epoch [24/60], Iter [50/157], Loss: total=2.496, reg=0.217, containing_obj=0.792, no_obj=0.733, cls=0.251

Epoch [24/60], Iter [100/157], Loss: total=2.478, reg=0.214, containing_obj=0.785, no_obj=0.740, cls=0.251

Epoch [24/60], Iter [150/157], Loss: total=2.479, reg=0.216, containing_obj=0.779, no_obj=0.740, cls=0.252

Starting epoch 25 / 60

Learning Rate for this epoch: 0.001

Epoch [25/60], Iter [50/157], Loss: total=2.543, reg=0.224, containing_obj=0.810, no_obj=0.721, cls=0.253

Epoch [25/60], Iter [100/157], Loss: total=2.525, reg=0.222, containing_obj=0.795, no_obj=0.725, cls=0.259

Epoch [25/60], Iter [150/157], Loss: total=2.486, reg=0.217, containing_obj=0.779, no_obj=0.731, cls=0.255

---Evaluate model on test samples---

100%|██████████| 1255/1255 [00:53<00:00, 23.30it/s]

```
---class aeroplane ap 0.5095303375452055---
---class bicycle ap 0.5636230268172828---
---class bird ap 0.4868602388369571---
---class boat ap 0.2557331174949898---
---class bottle ap 0.25031323081927626---
---class bus ap 0.43905409953157853---
---class car ap 0.5902147126052246---
---class cat ap 0.6138029832201106---
---class chair ap 0.30726784833897863---
---class cow ap 0.4888058132364483---
---class diningtable ap 0.2462738319881177---
---class dog ap 0.5798287063764304---
---class horse ap 0.5628228256651349---
---class motorbike ap 0.4903721358657399---
---class person ap 0.48483245619229187---
---class pottedplant ap 0.19852983374359182---
---class sheep ap 0.30695684433316023---
---class sofa ap 0.41033692023992163---
---class train ap 0.6903693063580909---
---class tvmonitor ap 0.5636455061022175---
---map 0.45195868876553746---
24 [0.5095303375452055, 0.5636230268172828, 0.4868602388369571, 0.255733117
4949898, 0.25031323081927626, 0.43905409953157853, 0.5902147126052246, 0.61
38029832201106, 0.30726784833897863, 0.4888058132364483, 0.246273831988117
7, 0.5798287063764304, 0.5628228256651349, 0.4903721358657399, 0.4848324561
9229187, 0.19852983374359182, 0.30695684433316023, 0.41033692023992163, 0.6
903693063580909, 0.5636455061022175]
```

Starting epoch 26 / 60

Learning Rate for this epoch: 0.001

Epoch [26/60], Iter [50/157], Loss: total=2.427, reg=0.216, containing_obj=0.767, no_obj=0.713, cls=0.225

Epoch [26/60], Iter [100/157], Loss: total=2.393, reg=0.212, containing_obj=0.738, no_obj=0.714, cls=0.240

Epoch [26/60], Iter [150/157], Loss: total=2.428, reg=0.213, containing_obj=0.764, no_obj=0.709, cls=0.247

Starting epoch 27 / 60

Learning Rate for this epoch: 0.001

Epoch [27/60], Iter [50/157], Loss: total=2.418, reg=0.212, containing_obj=0.756, no_obj=0.734, cls=0.236

Epoch [27/60], Iter [100/157], Loss: total=2.413, reg=0.212, containing_obj=0.749, no_obj=0.734, cls=0.236

Epoch [27/60], Iter [150/157], Loss: total=2.421, reg=0.212, containing_obj=0.758, no_obj=0.730, cls=0.237

Updating best val loss: 3.78541

Starting epoch 28 / 60

Learning Rate for this epoch: 0.001

Epoch [28/60], Iter [50/157], Loss: total=2.252, reg=0.197, containing_obj=0.700, no_obj=0.711, cls=0.212

Epoch [28/60], Iter [100/157], Loss: total=2.323, reg=0.203, containing_obj=0.723, no_obj=0.709, cls=0.228

Epoch [28/60], Iter [150/157], Loss: total=2.335, reg=0.202, containing_obj=0.733, no_obj=0.708, cls=0.237

Starting epoch 29 / 60

Learning Rate for this epoch: 0.001

Epoch [29/60], Iter [50/157], Loss: total=2.228, reg=0.198, containing_obj=0.687, no_obj=0.707, cls=0.198

Epoch [29/60], Iter [100/157], Loss: total=2.269, reg=0.200, containing_obj=0.698, no_obj=0.710, cls=0.216

Epoch [29/60], Iter [150/157], Loss: total=2.305, reg=0.201, containing_obj=0.730, no_obj=0.701, cls=0.217

Starting epoch 30 / 60

Learning Rate for this epoch: 0.001

Epoch [30/60], Iter [50/157], Loss: total=2.292, reg=0.201, containing_obj=0.733, no_obj=0.707, cls=0.198

Epoch [30/60], Iter [100/157], Loss: total=2.299, reg=0.202, containing_obj=0.722, no_obj=0.725, cls=0.206

Epoch [30/60], Iter [150/157], Loss: total=2.308, reg=0.203, containing_obj=0.728, no_obj=0.720, cls=0.207

---Evaluate model on test samples---

100%|██████████| 1255/1255 [00:54<00:00, 22.85it/s]


```
---class aeroplane ap 0.6177079387884008---
---class bicycle ap 0.5379795867432878---
---class bird ap 0.46769690909452794---
---class boat ap 0.3011821529571555---
---class bottle ap 0.16234119100641375---
---class bus ap 0.47799867421228814---
---class car ap 0.5934834294259874---
---class cat ap 0.6137310110142278---
---class chair ap 0.32715288188207126---
---class cow ap 0.4244127257773749---
---class diningtable ap 0.2430143393911509---
---class dog ap 0.5553825523010245---
---class horse ap 0.5397364128997498---
---class motorbike ap 0.3955670249565824---
---class person ap 0.42082889381186483---
---class pottedplant ap 0.14798643441824103---
---class sheep ap 0.2501305041035306---
---class sofa ap 0.4501045822944506---
---class train ap 0.6859155228309466---
---class tvmonitor ap 0.5844261135102953---
---map 0.43983894407097857---
29 [0.6177079387884008, 0.5379795867432878, 0.46769690909452794, 0.30118215
29571555, 0.16234119100641375, 0.47799867421228814, 0.5934834294259874, 0.6
137310110142278, 0.32715288188207126, 0.4244127257773749, 0.243014339391150
9, 0.5553825523010245, 0.5397364128997498, 0.3955670249565824, 0.4208288938
1186483, 0.14798643441824103, 0.2501305041035306, 0.4501045822944506, 0.685
9155228309466, 0.5844261135102953]
```

Starting epoch 31 / 60

Learning Rate for this epoch: 0.0001

Epoch [31/60], Iter [50/157], Loss: total=2.220, reg=0.200, containing_obj=0.694, no_obj=0.679, cls=0.186

Epoch [31/60], Iter [100/157], Loss: total=2.150, reg=0.191, containing_obj=0.669, no_obj=0.687, cls=0.184

Epoch [31/60], Iter [150/157], Loss: total=2.139, reg=0.191, containing_obj=0.664, no_obj=0.686, cls=0.180

Updating best val loss: 3.72051

Starting epoch 32 / 60

Learning Rate for this epoch: 0.0001

Epoch [32/60], Iter [50/157], Loss: total=2.087, reg=0.178, containing_obj=0.657, no_obj=0.684, cls=0.197

Epoch [32/60], Iter [100/157], Loss: total=2.082, reg=0.181, containing_obj=0.660, no_obj=0.680, cls=0.177

Epoch [32/60], Iter [150/157], Loss: total=2.057, reg=0.180, containing_obj=0.648, no_obj=0.679, cls=0.169

Updating best val loss: 3.70326

Starting epoch 33 / 60

Learning Rate for this epoch: 0.0001

Epoch [33/60], Iter [50/157], Loss: total=1.880, reg=0.163, containing_obj=0.586, no_obj=0.681, cls=0.140

Epoch [33/60], Iter [100/157], Loss: total=1.946, reg=0.169, containing_obj=0.614, no_obj=0.672, cls=0.152

Epoch [33/60], Iter [150/157], Loss: total=1.976, reg=0.171, containing_obj=0.631, no_obj=0.671, cls=0.156

Starting epoch 34 / 60

Learning Rate for this epoch: 0.0001

Epoch [34/60], Iter [50/157], Loss: total=2.038, reg=0.184, containing_obj=

```
0.632, no_obj=0.664, cls=0.155
Epoch [34/60], Iter [100/157], Loss: total=1.982, reg=0.177, containing_obj=
=0.608, no_obj=0.672, cls=0.152
Epoch [34/60], Iter [150/157], Loss: total=2.004, reg=0.179, containing_obj
=0.619, no_obj=0.671, cls=0.157
Updating best val loss: 3.69046
```

```
Starting epoch 35 / 60
Learning Rate for this epoch: 0.0001
Epoch [35/60], Iter [50/157], Loss: total=2.007, reg=0.175, containing_obj=
0.634, no_obj=0.662, cls=0.169
Epoch [35/60], Iter [100/157], Loss: total=2.001, reg=0.176, containing_obj
=0.634, no_obj=0.659, cls=0.155
Epoch [35/60], Iter [150/157], Loss: total=1.992, reg=0.175, containing_obj
=0.628, no_obj=0.673, cls=0.154
---Evaluate model on test samples---
```

```
100%|██████████| 1255/1255 [00:53<00:00, 23.33it/s]
```

```
---class aeroplane ap 0.5735311007734852---
---class bicycle ap 0.6210962880908787---
---class bird ap 0.5269692712287606---
---class boat ap 0.3458178432929257---
---class bottle ap 0.36489510912293355---
---class bus ap 0.5369657594161664---
---class car ap 0.6645655508709984---
---class cat ap 0.6992976093495711---
---class chair ap 0.4114211078109351---
---class cow ap 0.5709124250592041---
---class diningtable ap 0.3085704930706092---
---class dog ap 0.6133990462071444---
---class horse ap 0.6014821606144455---
---class motorbike ap 0.5392622162839699---
---class person ap 0.5323956861024063---
---class pottedplant ap 0.19597123388387344---
---class sheep ap 0.39345884292032496---
---class sofa ap 0.46830367303292814---
---class train ap 0.7588700558123438---
---class tvmonitor ap 0.6488596219999982---
---map 0.5188022547471951---
34 [0.5735311007734852, 0.6210962880908787, 0.5269692712287606, 0.345817843
2929257, 0.36489510912293355, 0.5369657594161664, 0.6645655508709984, 0.699
2976093495711, 0.4114211078109351, 0.5709124250592041, 0.3085704930706092,
0.6133990462071444, 0.6014821606144455, 0.5392622162839699, 0.5323956861024
063, 0.19597123388387344, 0.39345884292032496, 0.46830367303292814, 0.75887
00558123438, 0.6488596219999982]
Updating best val loss: 3.67718
```

Starting epoch 36 / 60

Learning Rate for this epoch: 0.0001

Epoch [36/60], Iter [50/157], Loss: total=2.056, reg=0.182, containing_obj=0.655, no_obj=0.668, cls=0.157

Epoch [36/60], Iter [100/157], Loss: total=1.935, reg=0.170, containing_obj=0.608, no_obj=0.669, cls=0.144

Epoch [36/60], Iter [150/157], Loss: total=1.952, reg=0.170, containing_obj=0.619, no_obj=0.669, cls=0.149

Starting epoch 37 / 60

Learning Rate for this epoch: 0.0001

Epoch [37/60], Iter [50/157], Loss: total=1.982, reg=0.173, containing_obj=0.631, no_obj=0.642, cls=0.166

Epoch [37/60], Iter [100/157], Loss: total=1.958, reg=0.172, containing_obj=0.623, no_obj=0.644, cls=0.155

Epoch [37/60], Iter [150/157], Loss: total=1.936, reg=0.170, containing_obj=0.612, no_obj=0.648, cls=0.151

Starting epoch 38 / 60

Learning Rate for this epoch: 0.0001

Epoch [38/60], Iter [50/157], Loss: total=1.874, reg=0.167, containing_obj=0.577, no_obj=0.662, cls=0.133

Epoch [38/60], Iter [100/157], Loss: total=1.932, reg=0.170, containing_obj=0.611, no_obj=0.653, cls=0.142

Epoch [38/60], Iter [150/157], Loss: total=1.918, reg=0.169, containing_obj=0.604, no_obj=0.659, cls=0.141

Starting epoch 39 / 60

Learning Rate for this epoch: 0.0001

Epoch [39/60], Iter [50/157], Loss: total=1.890, reg=0.168, containing_obj=0.587, no_obj=0.624, cls=0.149

Epoch [39/60], Iter [100/157], Loss: total=1.889, reg=0.165, containing_obj=0.598, no_obj=0.638, cls=0.147

Epoch [39/60], Iter [150/157], Loss: total=1.909, reg=0.167, containing_obj=0.606, no_obj=0.645, cls=0.144

Starting epoch 40 / 60

Learning Rate for this epoch: 0.0001

Epoch [40/60], Iter [50/157], Loss: total=1.953, reg=0.170, containing_obj=0.624, no_obj=0.662, cls=0.147

Epoch [40/60], Iter [100/157], Loss: total=1.921, reg=0.169, containing_obj=0.606, no_obj=0.656, cls=0.141

Epoch [40/60], Iter [150/157], Loss: total=1.936, reg=0.169, containing_obj=0.619, no_obj=0.654, cls=0.144

---Evaluate model on test samples---

100%|██████████| 1255/1255 [00:52<00:00, 23.77it/s]

```
---class aeroplane ap 0.6444464550004724---
---class bicycle ap 0.5831834467893946---
---class bird ap 0.5259217475487383---
---class boat ap 0.32998443027696894---
---class bottle ap 0.3412765027544556---
---class bus ap 0.5461452307215877---
---class car ap 0.6757382343206753---
---class cat ap 0.7028979822068162---
---class chair ap 0.3828901887518413---
---class cow ap 0.5628379951498034---
---class diningtable ap 0.2897252466657395---
---class dog ap 0.5910943408343462---
---class horse ap 0.6367554135924314---
---class motorbike ap 0.5617524659811451---
---class person ap 0.5515845948577743---
---class pottedplant ap 0.19956645934214837---
---class sheep ap 0.37321959274045835---
---class sofa ap 0.4576882910202789---
---class train ap 0.7276521436479862---
---class tvmonitor ap 0.6411873997723524---
---map 0.5162774080987708---
39 [0.6444464550004724, 0.5831834467893946, 0.5259217475487383, 0.329984430
27696894, 0.3412765027544556, 0.5461452307215877, 0.6757382343206753, 0.702
8979822068162, 0.3828901887518413, 0.5628379951498034, 0.2897252466657395,
0.5910943408343462, 0.6367554135924314, 0.5617524659811451, 0.5515845948577
743, 0.19956645934214837, 0.37321959274045835, 0.4576882910202789, 0.727652
1436479862, 0.6411873997723524]
```

Starting epoch 41 / 60

Learning Rate for this epoch: 1e-05

Epoch [41/60], Iter [50/157], Loss: total=1.810, reg=0.163, containing_obj=0.548, no_obj=0.632, cls=0.132

Epoch [41/60], Iter [100/157], Loss: total=1.896, reg=0.169, containing_obj=0.584, no_obj=0.646, cls=0.145

Epoch [41/60], Iter [150/157], Loss: total=1.906, reg=0.169, containing_obj=0.586, no_obj=0.655, cls=0.147

Starting epoch 42 / 60

Learning Rate for this epoch: 1e-05

Epoch [42/60], Iter [50/157], Loss: total=1.880, reg=0.166, containing_obj=0.599, no_obj=0.639, cls=0.132

Epoch [42/60], Iter [100/157], Loss: total=1.821, reg=0.158, containing_obj=0.567, no_obj=0.660, cls=0.133

Epoch [42/60], Iter [150/157], Loss: total=1.865, reg=0.162, containing_obj=0.582, no_obj=0.664, cls=0.140

Starting epoch 43 / 60

Learning Rate for this epoch: 1e-05

Epoch [43/60], Iter [50/157], Loss: total=1.905, reg=0.171, containing_obj=0.593, no_obj=0.643, cls=0.137

Epoch [43/60], Iter [100/157], Loss: total=1.932, reg=0.172, containing_obj=0.597, no_obj=0.654, cls=0.148

Epoch [43/60], Iter [150/157], Loss: total=1.902, reg=0.169, containing_obj=0.587, no_obj=0.656, cls=0.144

Starting epoch 44 / 60

Learning Rate for this epoch: 1e-05

Epoch [44/60], Iter [50/157], Loss: total=1.870, reg=0.164, containing_obj=0.574, no_obj=0.675, cls=0.138

Epoch [44/60], Iter [100/157], Loss: total=1.866, reg=0.164, containing_obj

```
=0.576, no_obj=0.668, cls=0.135  
Epoch [44/60], Iter [150/157], Loss: total=1.885, reg=0.166, containing_obj=  
=0.586, no_obj=0.661, cls=0.140
```

Starting epoch 45 / 60

Learning Rate for this epoch: 1e-05

```
Epoch [45/60], Iter [50/157], Loss: total=1.855, reg=0.164, containing_obj=  
0.571, no_obj=0.636, cls=0.144
```

```
Epoch [45/60], Iter [100/157], Loss: total=1.862, reg=0.166, containing_obj  
=0.573, no_obj=0.640, cls=0.141
```

```
Epoch [45/60], Iter [150/157], Loss: total=1.899, reg=0.167, containing_obj  
=0.595, no_obj=0.643, cls=0.145
```

---Evaluate model on test samples---

```
100%|██████████| 1255/1255 [00:53<00:00, 23.27it/s]
```

```
---class aeroplane ap 0.6423985603308207---
---class bicycle ap 0.5994162156318386---
---class bird ap 0.510893228544583---
---class boat ap 0.3547740564463797---
---class bottle ap 0.3662133504290982---
---class bus ap 0.5573304440323539---
---class car ap 0.6864546902074757---
---class cat ap 0.6935888604375255---
---class chair ap 0.3813829180557521---
---class cow ap 0.5771469556179596---
---class diningtable ap 0.30328886561979995---
---class dog ap 0.5886321666048049---
---class horse ap 0.6493583074908449---
---class motorbike ap 0.5389043242876863---
---class person ap 0.5505798661495862---
---class pottedplant ap 0.2040611578395672---
---class sheep ap 0.36273830124636575---
---class sofa ap 0.45887889422148365---
---class train ap 0.7446696497785146---
---class tvmonitor ap 0.6616024111588747---
---map 0.5216156612065657---
44 [0.6423985603308207, 0.5994162156318386, 0.510893228544583, 0.3547740564
463797, 0.3662133504290982, 0.5573304440323539, 0.6864546902074757, 0.69358
88604375255, 0.3813829180557521, 0.5771469556179596, 0.30328886561979995,
0.5886321666048049, 0.6493583074908449, 0.5389043242876863, 0.5505798661495
862, 0.2040611578395672, 0.36273830124636575, 0.45887889422148365, 0.744669
6497785146, 0.6616024111588747]
```

Starting epoch 46 / 60

Learning Rate for this epoch: 1e-05

Epoch [46/60], Iter [50/157], Loss: total=1.882, reg=0.169, containing_obj=0.576, no_obj=0.671, cls=0.128

Epoch [46/60], Iter [100/157], Loss: total=1.854, reg=0.165, containing_obj=0.566, no_obj=0.655, cls=0.133

Epoch [46/60], Iter [150/157], Loss: total=1.877, reg=0.167, containing_obj=0.579, no_obj=0.655, cls=0.135

Updating best val loss: 3.66392

Starting epoch 47 / 60

Learning Rate for this epoch: 1e-05

Epoch [47/60], Iter [50/157], Loss: total=1.869, reg=0.164, containing_obj=0.593, no_obj=0.657, cls=0.129

Epoch [47/60], Iter [100/157], Loss: total=1.857, reg=0.164, containing_obj=0.571, no_obj=0.656, cls=0.137

Epoch [47/60], Iter [150/157], Loss: total=1.913, reg=0.168, containing_obj=0.603, no_obj=0.654, cls=0.144

Updating best val loss: 3.66285

Starting epoch 48 / 60

Learning Rate for this epoch: 1e-05

Epoch [48/60], Iter [50/157], Loss: total=2.014, reg=0.180, containing_obj=0.643, no_obj=0.627, cls=0.159

Epoch [48/60], Iter [100/157], Loss: total=1.921, reg=0.170, containing_obj=0.608, no_obj=0.639, cls=0.142

Epoch [48/60], Iter [150/157], Loss: total=1.885, reg=0.167, containing_obj=0.589, no_obj=0.643, cls=0.139

Updating best val loss: 3.64994

Starting epoch 49 / 60

Learning Rate for this epoch: 1e-05

```
Epoch [49/60], Iter [50/157], Loss: total=1.882, reg=0.168, containing_obj=
0.579, no_obj=0.626, cls=0.151
Epoch [49/60], Iter [100/157], Loss: total=1.866, reg=0.165, containing_obj
=0.581, no_obj=0.633, cls=0.145
Epoch [49/60], Iter [150/157], Loss: total=1.889, reg=0.166, containing_obj
=0.592, no_obj=0.643, cls=0.143
```

Starting epoch 50 / 60

Learning Rate for this epoch: 1e-05

```
Epoch [50/60], Iter [50/157], Loss: total=1.874, reg=0.168, containing_obj=
0.561, no_obj=0.642, cls=0.153
```

```
Epoch [50/60], Iter [100/157], Loss: total=1.943, reg=0.175, containing_obj
=0.594, no_obj=0.648, cls=0.151
```

```
Epoch [50/60], Iter [150/157], Loss: total=1.900, reg=0.168, containing_obj
=0.588, no_obj=0.652, cls=0.145
```

---Evaluate model on test samples---

```
100%|██████████| 1255/1255 [00:53<00:00, 23.48it/s]
```



```
---class aeroplane ap 0.6384338474284622---
---class bicycle ap 0.6136327785694097---
---class bird ap 0.5259013196541708---
---class boat ap 0.3263655745083587---
---class bottle ap 0.3650259569367671---
---class bus ap 0.5602035724935875---
---class car ap 0.6868775945441385---
---class cat ap 0.707258812199595---
---class chair ap 0.392217815218111---
---class cow ap 0.5681528308418516---
---class diningtable ap 0.3228357094370166---
---class dog ap 0.6151056262669146---
---class horse ap 0.6326052045566848---
---class motorbike ap 0.5674421802046101---
---class person ap 0.5511173385203563---
---class pottedplant ap 0.20657225054030856---
---class sheep ap 0.3829913763845815---
---class sofa ap 0.46942110099394513---
---class train ap 0.732203683885809---
---class tvmonitor ap 0.661408288710597---
---map 0.5262886430947638---
49 [0.6384338474284622, 0.6136327785694097, 0.5259013196541708, 0.326365574
5083587, 0.3650259569367671, 0.5602035724935875, 0.6868775945441385, 0.7072
58812199595, 0.392217815218111, 0.5681528308418516, 0.3228357094370166, 0.6
151056262669146, 0.6326052045566848, 0.5674421802046101, 0.551117338520356
3, 0.20657225054030856, 0.3829913763845815, 0.46942110099394513, 0.73220368
3885809, 0.661408288710597]
```

Starting epoch 51 / 60

Learning Rate for this epoch: 1e-05

Epoch [51/60], Iter [50/157], Loss: total=1.807, reg=0.154, containing_obj=0.575, no_obj=0.651, cls=0.137

Epoch [51/60], Iter [100/157], Loss: total=1.858, reg=0.162, containing_obj=0.583, no_obj=0.651, cls=0.138

Epoch [51/60], Iter [150/157], Loss: total=1.865, reg=0.162, containing_obj=0.585, no_obj=0.648, cls=0.145

Starting epoch 52 / 60

Learning Rate for this epoch: 1e-05

Epoch [52/60], Iter [50/157], Loss: total=1.925, reg=0.169, containing_obj=0.608, no_obj=0.644, cls=0.149

Epoch [52/60], Iter [100/157], Loss: total=1.874, reg=0.165, containing_obj=0.584, no_obj=0.654, cls=0.139

Epoch [52/60], Iter [150/157], Loss: total=1.848, reg=0.162, containing_obj=0.574, no_obj=0.656, cls=0.136

Starting epoch 53 / 60

Learning Rate for this epoch: 1e-05

Epoch [53/60], Iter [50/157], Loss: total=1.865, reg=0.164, containing_obj=0.582, no_obj=0.649, cls=0.137

Epoch [53/60], Iter [100/157], Loss: total=1.867, reg=0.163, containing_obj=0.584, no_obj=0.650, cls=0.141

Epoch [53/60], Iter [150/157], Loss: total=1.874, reg=0.164, containing_obj=0.587, no_obj=0.650, cls=0.144

Starting epoch 54 / 60

Learning Rate for this epoch: 1e-05

Epoch [54/60], Iter [50/157], Loss: total=1.818, reg=0.159, containing_obj=0.555, no_obj=0.670, cls=0.135

Epoch [54/60], Iter [100/157], Loss: total=1.839, reg=0.159, containing_obj

```
=0.571, no_obj=0.662, cls=0.142
Epoch [54/60], Iter [150/157], Loss: total=1.861, reg=0.162, containing_obj
=0.587, no_obj=0.655, cls=0.137
```

Starting epoch 55 / 60

Learning Rate for this epoch: 1e-05

```
Epoch [55/60], Iter [50/157], Loss: total=1.798, reg=0.156, containing_obj=
0.551, no_obj=0.654, cls=0.140
```

```
Epoch [55/60], Iter [100/157], Loss: total=1.849, reg=0.161, containing_obj
=0.580, no_obj=0.647, cls=0.140
```

```
Epoch [55/60], Iter [150/157], Loss: total=1.854, reg=0.164, containing_obj
=0.574, no_obj=0.650, cls=0.135
```

---Evaluate model on test samples---

```
100%|██████████| 1255/1255 [00:54<00:00, 22.93it/s]
```

```
---class aeroplane ap 0.6335184556932523---
```

```
---class bicycle ap 0.5917583435812003---
```

```
---class bird ap 0.5234884068276374---
```

```
---class boat ap 0.3494766127912661---
```

```
---class bottle ap 0.35724440816858893---
```

```
---class bus ap 0.5255734594308085---
```

```
---class car ap 0.6811668767126993---
```

```
---class cat ap 0.7246887851801376---
```

```
---class chair ap 0.3931015015042486---
```

```
---class cow ap 0.5742979496715412---
```

```
---class diningtable ap 0.3277197587933954---
```

```
---class dog ap 0.5878976421260815---
```

```
---class horse ap 0.676276914334822---
```

```
---class motorbike ap 0.5410341101907261---
```

```
---class person ap 0.5537199911590934---
```

```
---class pottedplant ap 0.19078454041899562---
```

```
---class sheep ap 0.3585872832407281---
```

```
---class sofa ap 0.4392835229624705---
```

```
---class train ap 0.7161043768417188---
```

```
---class tvmonitor ap 0.6511845295439105---
```

```
---map 0.5198453734586661---
```

```
54 [0.6335184556932523, 0.5917583435812003, 0.5234884068276374, 0.349476612
7912661, 0.35724440816858893, 0.5255734594308085, 0.6811668767126993, 0.724
6887851801376, 0.3931015015042486, 0.5742979496715412, 0.3277197587933954,
0.5878976421260815, 0.676276914334822, 0.5410341101907261, 0.55371999115909
34, 0.19078454041899562, 0.3585872832407281, 0.4392835229624705, 0.71610437
68417188, 0.6511845295439105]
```

Starting epoch 56 / 60

Learning Rate for this epoch: 1e-05

```
Epoch [56/60], Iter [50/157], Loss: total=1.964, reg=0.172, containing_obj=
0.639, no_obj=0.640, cls=0.143
```

```
Epoch [56/60], Iter [100/157], Loss: total=1.860, reg=0.162, containing_obj
=0.592, no_obj=0.652, cls=0.132
```

```
Epoch [56/60], Iter [150/157], Loss: total=1.849, reg=0.161, containing_obj
=0.581, no_obj=0.656, cls=0.136
```

Starting epoch 57 / 60

Learning Rate for this epoch: 1e-05

```
Epoch [57/60], Iter [50/157], Loss: total=1.817, reg=0.160, containing_obj=
0.546, no_obj=0.646, cls=0.146
```

```
Epoch [57/60], Iter [100/157], Loss: total=1.895, reg=0.167, containing_obj
=0.581, no_obj=0.649, cls=0.155
```

View example predictions

```
In [ ]: net.eval()

# select random image from val set
image_name = random.choice(val_dataset.fnames)
image = cv2.imread(os.path.join(file_root_val, image_name))
image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

print('predicting...')
result = predict_image(net, image_name, root_img_directory=file_root_val)
for left_up, right_bottom, class_name, _, prob in result:
    color = COLORS[VOC_CLASSES.index(class_name)]
    cv2.rectangle(image, left_up, right_bottom, color, 2)
    label = class_name + str(round(prob, 2))
    text_size, baseline = cv2.getTextSize(label, cv2.FONT_HERSHEY_SIMPLEX, 0.8, 1)
    p1 = (left_up[0], left_up[1] - text_size[1])
    cv2.rectangle(image, (p1[0] - 2 // 2, p1[1] - 2 - baseline), (p1[0] + text_size[0] + 2, p1[1]),
                  color, -1)
    cv2.putText(image, label, (p1[0], p1[1] + baseline), cv2.FONT_HERSHEY_SIMPLEX, 0.8, color)

plt.figure(figsize = (15,15))
plt.imshow(image)
```

Kaggle submission (85%)

Predict Result

Predict the results based on testing set. Upload to [Kaggle](#).

How to upload

1. Click the folder icon in the left hand side of Colab.
2. Right click "result.csv". Select "Download"
3. To kaggle. Click "Submit Predictions"
4. Upload the result.csv
5. System will automaticlaly calculate the accuracy of 50% dataset and publish this result to leaderboard.

預測 test 並將結果上傳至Kaggle。 [連結](#)

執行完畢此區的程式碼後，會將 test 預測完的結果存下來。

上傳流程

1. 點選左側選單最下方的資料夾圖示
2. 右鍵「result.csv」
3. 點選「Download」
4. 至連結網頁點選「Submit Predictions」
5. 將剛剛下載的檔案上傳
6. 系統會計算並公布其中50%資料的正確率

```
In [ ]: root_test = 'data/VOCdevkit_2007_2/VOC2007test/JPEGImages/'
file_test = 'data/voc2007test.txt'
```

By using the `test_evaluate` function, you will obtain predictions for each image.

```
In [ ]: preds_submission = test_evaluate(net, test_dataset_file=file_test, img_root=
```

The `write_csv` function will use `preds_submission` to write into a CSV file called 'result.csv'.

```
In [ ]: write_csv(preds_submission)
```

```
In [ ]: from google.colab import drive
drive.mount('/content/drive')
```

Report (15%)

In your report, please include:

- a. A brief discussion on your implementation.
- b. Report the best train and validation accuracy in all of your experiments and discuss any strategies or tricks you've employed.
- c. Report the results for extra credits and also provide a discussion, if any.

Extra Credit (15%)

- Pick a fun video like [this one](#), run your detector on it (a subset of frames would be OK), and produce a video showing your results.
- Try to replace the provided pre-trained network with a different one and train with the YOLO loss on top to attempt to get better accuracy.
- Or any other methods that you try to improve the performance.

MIS 583 Assignment 5: YOLO Object Detection on PASCAL VOC

Before we start, please put your name and SID in following format:

: LASTNAME Firstname, ?00000000 // e.g.) 李晨愷 M114020035

Your Answer:

Hi I'm 游雅淇, B104020012.

Google Colab Setup

Next we need to run a few commands to set up our environment on Google Colab. If you are running this notebook on a local machine you can skip this section.

Run the following cell to mount your Google Drive. Follow the link, sign in to your Google account (the same account you used to store this notebook!) and copy the authorization code into the text box that appears below.

```
In [ ]: from google.colab import drive
drive.mount('/content/drive', force_remount=True)
```

Mounted at /content/drive

How to Get Data

請先到共用雲端硬碟將檔案 `VOCdevkit_2007.zip`，建立捷徑到自己的雲端硬碟中。

操作步驟

1. 點開雲端[連結](#)
2. 點選右上角「新增雲端硬碟捷徑」
3. 點選「我的雲端硬碟」
4. 點選「新增捷徑」

完成以上流程會在你的雲端硬碟中建立一個檔案的捷徑，接著我們在colab中取得權限即可使用。

Unzip Data

解壓縮 `VOCdevkit_2007.zip`

- `VOC2007` : 包含了train/val的所有圖片
- `VOC2007test` : 包含了test的所有圖片

其中 `train` 的圖片 3756 張，`val` 的圖片 1255 張，`test` 的圖片 4950 張。

注意: 若有另外設定存放在雲端硬碟中的路徑，請記得本處路徑也須做更動。

Notice: Please put "VOCdevkit_2007" folder under data folder.

```
In [ ]: !unzip -qq ./drive/MyDrive/DeepLearning/A5/data/VOCdevkit_2007.zip
```

```
In [ ]: %cd ./drive/MyDrive/DeepLearning/A5
/content/drive/MyDrive/DeepLearning/A5
```

Import package

```
In [ ]: import os
import random

import cv2
import numpy as np

import csv

import torch
from torch.utils.data import DataLoader
from torchvision import models

from src.resnet_yolo import resnet50
from src.densenet_yolo import densenet121
from yolo_loss import YoloLoss
from src.dataset import VocDetectorDataset
from src.eval_voc import evaluate, test_evaluate
from src.predict import predict_image
from src.config import VOC_CLASSES, COLORS
from kaggle_submission import write_csv

import matplotlib.pyplot as plt
import collections

%matplotlib inline
%load_ext autoreload
%autoreload 2
```

Initialization

```
In [ ]: device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
```

```
In [ ]: # YOLO network hyperparameters
B = 2 # number of bounding box predictions per cell
S = 14 # width/height of network output grid (larger than 7x7 from paper s
```

To implement Yolo we will rely on a pretrained classifier as the backbone for our detection network. PyTorch offers a variety of models which are pretrained on ImageNet in the `torchvision.models` package. In particular, we will use the ResNet50 architecture as a base for our detector. This is different from the base architecture in the Yolo paper and also results in a different output grid size (14x14 instead of 7x7).

Models are typically pretrained on ImageNet since the dataset is very large (> 1 million images) and widely used. The pretrained model provides a very useful weight initialization for our detector, so that the network is able to learn quickly and effectively.

Resnet50

```
In [ ]: load_network_path = None #'checkpoints/best_detector.pth'
pretrained = True

# use to load a previously trained network
if load_network_path is not None:
    print('Loading saved network from {}'.format(load_network_path))
    net = resnet50().to(device)
    net.load_state_dict(torch.load(load_network_path))
else:
    print('Load pre-trained model')
    net = resnet50(pretrained=pretrained).to(device)
```

Load pre-trained model

```
/usr/local/lib/python3.10/dist-packages/torchvision/models/_utils.py:208: U
serWarning: The parameter 'pretrained' is deprecated since 0.13 and may be
removed in the future, please use 'weights' instead.
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/torchvision/models/_utils.py:223: U
serWarning: Arguments other than a weight enum or `None` for 'weights' are
deprecated since 0.13 and may be removed in the future. The current behavio
r is equivalent to passing `weights=ResNet50_Weights.IMAGENET1K_V1`. You ca
n also use `weights=ResNet50_Weights.DEFAULT` to get the most up-to-date we
ights.
  warnings.warn(msg)
Downloading: "https://download.pytorch.org/models/resnet50-0676ba61.pth" to
/root/.cache/torch/hub/checkpoints/resnet50-0676ba61.pth
100%|██████████| 97.8M/97.8M [00:00<00:00, 165MB/s]
```

DenseNet121

```
In [ ]: # Densenet121
load_network_path = None #'checkpoints/best_detector.pth'
pretrained = True

# use to load a previously trained network
if load_network_path is not None:
    print('Loading saved network from {}'.format(load_network_path))
    net = densenet121(S=S).to(device)
    net.load_state_dict(torch.load(load_network_path))
else:
    print('Load pre-trained model')
    net = densenet121(pretrained=pretrained, S=S).to(device)
```

```
In [ ]: learning_rate = 0.001
num_epochs = 50
batch_size = 24

# Yolo loss component coefficients (as given in Yolo v1 paper)
lambda_coord = 5
lambda_noobj = 0.5
```

Reading Pascal Data

The train dataset loader also using a variety of data augmentation techniques including random shift, scaling, crop, and flips. Data augmentation is slightly more complicated for

detection datasets since the bounding box annotations must be kept consistent throughout the transformations.

Since the output of the detector network we train is an $S \times S \times (B \times 5 + C)$, we use an encoder to convert the original bounding box coordinates into relative grid bounding box coordinates corresponding to the expected output. We also use a decoder which allows us to convert the opposite direction into image coordinate bounding boxes.

Notice: Please put "VOCdevkit_2007" folder under data folder.

```
In [ ]: file_root_train = 'data/VOCdevkit_2007/VOC2007/JPEGImages/'
        annotation_file_train = 'data/voc2007train.txt'

        train_dataset = VocDetectorDataset(root_img_dir=file_root_train, dataset_file=annotation_file_train)
        train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True, num_workers=4)
        print('Loaded %d train images' % len(train_dataset))
```

Initializing dataset
Loaded 3756 train images

```
In [ ]: file_root_val = 'data/VOCdevkit_2007/VOC2007/JPEGImages/'
        annotation_file_val = 'data/voc2007val.txt'

        val_dataset = VocDetectorDataset(root_img_dir=file_root_val, dataset_file=annotation_file_val)
        val_loader = DataLoader(val_dataset, batch_size=batch_size, shuffle=False, num_workers=4)
        print('Loaded %d val images' % len(val_dataset))
```

Initializing dataset
Loaded 1255 val images

```
In [ ]: data = train_dataset[0]
```

Set up training tools

```
In [ ]: criterion = YoloLoss(S, B, lambda_coord, lambda_noobj)
        optimizer = torch.optim.SGD(net.parameters(), lr=learning_rate, momentum=0.9)
```

Train detector

```
In [ ]: best_val_loss = np.inf
        learning_rate = 1e-3
        for epoch in range(num_epochs):
            net.train()

            # Update learning rate late in training
            if epoch == 30 or epoch == 40 or epoch == 50 :
                learning_rate /= 10.0

            for param_group in optimizer.param_groups:
                param_group['lr'] = learning_rate

            print('\n\nStarting epoch %d / %d' % (epoch + 1, num_epochs))
            print('Learning Rate for this epoch: {}'.format(learning_rate))

            total_loss = collections.defaultdict(int)

            for i, data in enumerate(train_loader):
                data = (item.to(device) for item in data)
```



```

images, target_boxes, target_cls, has_object_map = data
pred = net(images)
loss_dict = criterion(pred, target_boxes, target_cls, has_object_map)
for key in loss_dict:
    total_loss[key] += loss_dict[key].item()

optimizer.zero_grad()
loss_dict['total_loss'].backward()
optimizer.step()

if (i+1) % 50 == 0:
    outstring = 'Epoch [%d/%d], Iter [%d/%d], Loss: ' % ((epoch+1), i+1, i+1, i+1)
    outstring += ', '.join( "%s=%.3f" % (key[:5], val / (i+1)) for key, val in loss_dict.items())
    print(outstring)

# evaluate the network on the val data
if (epoch + 1) % 5 == 0:
    val_aps = evaluate(net, val_dataset_file=annotation_file_val, img_root=val_root)
    print(epoch, val_aps)
with torch.no_grad():
    val_loss = 0.0
    net.eval()
    for i, data in enumerate(val_loader):
        data = (item.to(device) for item in data)
        images, target_boxes, target_cls, has_object_map = data

        pred = net(images)
        loss_dict = criterion(pred, target_boxes, target_cls, has_object_map)
        val_loss += loss_dict['total_loss'].item()
    val_loss /= len(val_loader)

if best_val_loss > val_loss:
    best_val_loss = val_loss
    print('Updating best val loss: %.5f' % best_val_loss)
    torch.save(net.state_dict(), 'checkpoints/best_detector.pth')

if (epoch+1) in [5, 10, 20, 30, 40, 50, 60, 70, 80, 90, 100]:
    torch.save(net.state_dict(), 'checkpoints/detector_epoch_%d.pth' % (epoch+1))

torch.save(net.state_dict(), 'checkpoints/detector.pth')

```

Starting epoch 1 / 50
Learning Rate for this epoch: 0.001
Epoch [1/50], Iter [50/157], Loss: total=24.103, reg=0.758, containing_obj=1.414, no_obj=25.251, cls=6.271
Epoch [1/50], Iter [100/157], Loss: total=16.279, reg=0.631, containing_obj=1.722, no_obj=13.118, cls=4.844
Epoch [1/50], Iter [150/157], Loss: total=13.335, reg=0.576, containing_obj=1.826, no_obj=8.995, cls=4.133
Updating best val loss: 7.04412

Starting epoch 2 / 50
Learning Rate for this epoch: 0.001
Epoch [2/50], Iter [50/157], Loss: total=6.573, reg=0.428, containing_obj=1.887, no_obj=0.681, cls=2.204
Epoch [2/50], Iter [100/157], Loss: total=6.342, reg=0.420, containing_obj=1.799, no_obj=0.725, cls=2.078
Epoch [2/50], Iter [150/157], Loss: total=6.245, reg=0.419, containing_obj=1.751, no_obj=0.761, cls=2.016
Updating best val loss: 6.16320

Starting epoch 3 / 50
Learning Rate for this epoch: 0.001
Epoch [3/50], Iter [50/157], Loss: total=5.470, reg=0.389, containing_obj=1.523, no_obj=0.811, cls=1.595
Epoch [3/50], Iter [100/157], Loss: total=5.573, reg=0.401, containing_obj=1.570, no_obj=0.801, cls=1.599
Epoch [3/50], Iter [150/157], Loss: total=5.488, reg=0.396, containing_obj=1.543, no_obj=0.808, cls=1.564
Updating best val loss: 5.70232

Starting epoch 4 / 50
Learning Rate for this epoch: 0.001
Epoch [4/50], Iter [50/157], Loss: total=5.366, reg=0.391, containing_obj=1.528, no_obj=0.841, cls=1.465
Epoch [4/50], Iter [100/157], Loss: total=5.109, reg=0.378, containing_obj=1.467, no_obj=0.811, cls=1.348
Epoch [4/50], Iter [150/157], Loss: total=5.031, reg=0.373, containing_obj=1.454, no_obj=0.795, cls=1.315
Updating best val loss: 5.27348

Starting epoch 5 / 50
Learning Rate for this epoch: 0.001
Epoch [5/50], Iter [50/157], Loss: total=4.898, reg=0.375, containing_obj=1.408, no_obj=0.790, cls=1.219
Epoch [5/50], Iter [100/157], Loss: total=4.730, reg=0.361, containing_obj=1.393, no_obj=0.767, cls=1.148
Epoch [5/50], Iter [150/157], Loss: total=4.690, reg=0.361, containing_obj=1.385, no_obj=0.764, cls=1.119
---Evaluate model on test samples---

100%|██████████| 1255/1255 [00:50<00:00, 24.82it/s]

```
---class aeroplane ap 0.23404448973733707---
---class bicycle ap 0.2674225959214281---
---class bird ap 0.3158070274510974---
---class boat ap 0.08444207196280162---
---class bottle ap 0.023050210899984915---
---class bus ap 0.07844155844155844---
---class car ap 0.33017389990423207---
---class cat ap 0.32077132121363017---
---class chair ap 0.15445498904886262---
---class cow ap 0.0608330288637194---
---class diningtable ap 0.031746031746031744---
---class dog ap 0.13054914156015418---
---class horse ap 0.31551656224285335---
---class motorbike ap 0.050271739130434784---
---class person ap 0.320592429697816---
---class pottedplant ap 0.0400545531069177---
---class sheep ap 0.12137958935623835---
---class sofa ap 0.041666666666666664---
---class train ap 0.30149593114709394---
---class tvmonitor ap 0.2850767792967337---
---map 0.17538953086977963---
4 [0.23404448973733707, 0.2674225959214281, 0.3158070274510974, 0.084442071
96280162, 0.023050210899984915, 0.07844155844155844, 0.33017389990423207,
0.32077132121363017, 0.15445498904886262, 0.0608330288637194, 0.03174603174
6031744, 0.13054914156015418, 0.31551656224285335, 0.050271739130434784, 0.
320592429697816, 0.0400545531069177, 0.12137958935623835, 0.041666666666666
664, 0.30149593114709394, 0.2850767792967337]
Updating best val loss: 5.04816
```

Starting epoch 6 / 50

Learning Rate for this epoch: 0.001

Epoch [6/50], Iter [50/157], Loss: total=4.375, reg=0.339, containing_obj=1.308, no_obj=0.767, cls=0.988

Epoch [6/50], Iter [100/157], Loss: total=4.360, reg=0.337, containing_obj=1.318, no_obj=0.774, cls=0.971

Epoch [6/50], Iter [150/157], Loss: total=4.365, reg=0.342, containing_obj=1.314, no_obj=0.775, cls=0.955

Updating best val loss: 4.79899

Starting epoch 7 / 50

Learning Rate for this epoch: 0.001

Epoch [7/50], Iter [50/157], Loss: total=4.086, reg=0.324, containing_obj=1.239, no_obj=0.784, cls=0.837

Epoch [7/50], Iter [100/157], Loss: total=4.106, reg=0.326, containing_obj=1.243, no_obj=0.774, cls=0.845

Epoch [7/50], Iter [150/157], Loss: total=4.115, reg=0.329, containing_obj=1.248, no_obj=0.765, cls=0.839

Updating best val loss: 4.72316

Starting epoch 8 / 50

Learning Rate for this epoch: 0.001

Epoch [8/50], Iter [50/157], Loss: total=4.122, reg=0.329, containing_obj=1.268, no_obj=0.786, cls=0.814

Epoch [8/50], Iter [100/157], Loss: total=3.999, reg=0.320, containing_obj=1.224, no_obj=0.779, cls=0.786

Epoch [8/50], Iter [150/157], Loss: total=3.998, reg=0.322, containing_obj=1.218, no_obj=0.777, cls=0.783

Updating best val loss: 4.62204

Starting epoch 9 / 50

```
Learning Rate for this epoch: 0.001
Epoch [9/50], Iter [50/157], Loss: total=3.807, reg=0.313, containing_obj=
1.173, no_obj=0.787, cls=0.673
Epoch [9/50], Iter [100/157], Loss: total=3.846, reg=0.317, containing_obj=
1.185, no_obj=0.791, cls=0.680
Epoch [9/50], Iter [150/157], Loss: total=3.824, reg=0.313, containing_obj=
1.181, no_obj=0.785, cls=0.683
Updating best val loss: 4.43981
```

```
Starting epoch 10 / 50
Learning Rate for this epoch: 0.001
Epoch [10/50], Iter [50/157], Loss: total=3.753, reg=0.314, containing_obj=
1.177, no_obj=0.765, cls=0.625
Epoch [10/50], Iter [100/157], Loss: total=3.686, reg=0.305, containing_obj
=1.146, no_obj=0.777, cls=0.627
Epoch [10/50], Iter [150/157], Loss: total=3.695, reg=0.305, containing_obj
=1.154, no_obj=0.772, cls=0.628
---Evaluate model on test samples---
100%|██████████| 1255/1255 [00:44<00:00, 28.33it/s]
```

```
---class aeroplane ap 0.41599905429399675---
---class bicycle ap 0.3650426306577636---
---class bird ap 0.43968664510398064---
---class boat ap 0.23028502981241755---
---class bottle ap 0.07277936130933406---
---class bus ap 0.36589732992018886---
---class car ap 0.5150837674287456---
---class cat ap 0.6245462882596137---
---class chair ap 0.20866950602831133---
---class cow ap 0.20186145807263822---
---class diningtable ap 0.17838287047196416---
---class dog ap 0.44923107279254954---
---class horse ap 0.46145730771109944---
---class motorbike ap 0.3802583905415713---
---class person ap 0.43237009481686006---
---class pottedplant ap 0.120868787584249---
---class sheep ap 0.18854576502493667---
---class sofa ap 0.31401575662560055---
---class train ap 0.462676251948982---
---class tvmonitor ap 0.4626558406795302---
---map 0.34451566045421667---
9 [0.41599905429399675, 0.3650426306577636, 0.43968664510398064, 0.23028502
981241755, 0.07277936130933406, 0.36589732992018886, 0.5150837674287456, 0.
6245462882596137, 0.20866950602831133, 0.20186145807263822, 0.1783828704719
6416, 0.44923107279254954, 0.46145730771109944, 0.3802583905415713, 0.43237
009481686006, 0.120868787584249, 0.18854576502493667, 0.31401575662560055,
0.462676251948982, 0.4626558406795302]
Updating best val loss: 4.43909
```

```
Starting epoch 11 / 50
Learning Rate for this epoch: 0.001
Epoch [11/50], Iter [50/157], Loss: total=3.413, reg=0.281, containing_obj=
1.069, no_obj=0.735, cls=0.571
Epoch [11/50], Iter [100/157], Loss: total=3.528, reg=0.292, containing_obj
=1.102, no_obj=0.765, cls=0.582
Epoch [11/50], Iter [150/157], Loss: total=3.560, reg=0.295, containing_obj
=1.111, no_obj=0.771, cls=0.589
Updating best val loss: 4.36562
```

```
Starting epoch 12 / 50
Learning Rate for this epoch: 0.001
Epoch [12/50], Iter [50/157], Loss: total=3.348, reg=0.281, containing_obj=
1.048, no_obj=0.776, cls=0.506
Epoch [12/50], Iter [100/157], Loss: total=3.423, reg=0.286, containing_obj
=1.071, no_obj=0.784, cls=0.531
Epoch [12/50], Iter [150/157], Loss: total=3.432, reg=0.286, containing_obj
=1.073, no_obj=0.780, cls=0.538
Updating best val loss: 4.25210
```

```
Starting epoch 13 / 50
Learning Rate for this epoch: 0.001
Epoch [13/50], Iter [50/157], Loss: total=3.374, reg=0.280, containing_obj=
1.089, no_obj=0.724, cls=0.522
Epoch [13/50], Iter [100/157], Loss: total=3.324, reg=0.279, containing_obj
=1.061, no_obj=0.751, cls=0.494
Epoch [13/50], Iter [150/157], Loss: total=3.302, reg=0.277, containing_obj
=1.049, no_obj=0.763, cls=0.489
Updating best val loss: 4.13852
```

```
Starting epoch 14 / 50
```

```
Learning Rate for this epoch: 0.001
Epoch [14/50], Iter [50/157], Loss: total=3.192, reg=0.269, containing_obj=
0.996, no_obj=0.765, cls=0.468
Epoch [14/50], Iter [100/157], Loss: total=3.203, reg=0.270, containing_obj
=1.009, no_obj=0.767, cls=0.459
Epoch [14/50], Iter [150/157], Loss: total=3.226, reg=0.272, containing_obj
=1.019, no_obj=0.777, cls=0.461
```

Starting epoch 15 / 50

```
Learning Rate for this epoch: 0.001
Epoch [15/50], Iter [50/157], Loss: total=3.094, reg=0.261, containing_obj=
0.988, no_obj=0.757, cls=0.423
Epoch [15/50], Iter [100/157], Loss: total=3.173, reg=0.268, containing_obj
=1.008, no_obj=0.762, cls=0.445
Epoch [15/50], Iter [150/157], Loss: total=3.178, reg=0.270, containing_obj
=1.008, no_obj=0.769, cls=0.437
---Evaluate model on test samples---
```

```
100%|██████████| 1255/1255 [00:43<00:00, 28.92it/s]
```

```
---class aeroplane ap 0.3344076871015168---
---class bicycle ap 0.5425683158298162---
---class bird ap 0.5413258457213874---
---class boat ap 0.1748446140741921---
---class bottle ap 0.08956548758759073---
---class bus ap 0.44675533773595416---
---class car ap 0.5692992441660448---
---class cat ap 0.6372950048910108---
---class chair ap 0.21840869015689787---
---class cow ap 0.43179382275070716---
---class diningtable ap 0.2503938327467739---
---class dog ap 0.49069522054881254---
---class horse ap 0.5463111492013286---
---class motorbike ap 0.4347249289374292---
---class person ap 0.46142858756641214---
---class pottedplant ap 0.11552953546726205---
---class sheep ap 0.266905456607383---
---class sofa ap 0.3241537033666099---
---class train ap 0.5533180439960101---
---class tvmonitor ap 0.48334994517193763---
---map 0.39565372268125387---
14 [0.3344076871015168, 0.5425683158298162, 0.5413258457213874, 0.174844614
0741921, 0.08956548758759073, 0.44675533773595416, 0.5692992441660448, 0.63
72950048910108, 0.21840869015689787, 0.43179382275070716, 0.250393832746773
9, 0.49069522054881254, 0.5463111492013286, 0.4347249289374292, 0.461428587
56641214, 0.11552953546726205, 0.266905456607383, 0.3241537033666099, 0.553
3180439960101, 0.48334994517193763]
```

Starting epoch 16 / 50

Learning Rate for this epoch: 0.001

Epoch [16/50], Iter [50/157], Loss: total=3.134, reg=0.263, containing_obj=1.001, no_obj=0.761, cls=0.437

Epoch [16/50], Iter [100/157], Loss: total=3.063, reg=0.260, containing_obj=0.961, no_obj=0.769, cls=0.417

Epoch [16/50], Iter [150/157], Loss: total=3.090, reg=0.263, containing_obj=0.973, no_obj=0.768, cls=0.416

Updating best val loss: 4.04159

Starting epoch 17 / 50

Learning Rate for this epoch: 0.001

Epoch [17/50], Iter [50/157], Loss: total=2.996, reg=0.252, containing_obj=0.954, no_obj=0.767, cls=0.396

Epoch [17/50], Iter [100/157], Loss: total=2.925, reg=0.250, containing_obj=0.917, no_obj=0.777, cls=0.372

Epoch [17/50], Iter [150/157], Loss: total=2.967, reg=0.252, containing_obj=0.941, no_obj=0.769, cls=0.383

Starting epoch 18 / 50

Learning Rate for this epoch: 0.001

Epoch [18/50], Iter [50/157], Loss: total=2.873, reg=0.244, containing_obj=0.926, no_obj=0.763, cls=0.348

Epoch [18/50], Iter [100/157], Loss: total=2.915, reg=0.248, containing_obj=0.919, no_obj=0.768, cls=0.374

Epoch [18/50], Iter [150/157], Loss: total=2.909, reg=0.247, containing_obj=0.919, no_obj=0.762, cls=0.375

Updating best val loss: 3.99177

Starting epoch 19 / 50

Learning Rate for this epoch: 0.001

Epoch [19/50], Iter [50/157], Loss: total=2.853, reg=0.244, containing_obj=

```
0.905, no_obj=0.750, cls=0.353
Epoch [19/50], Iter [100/157], Loss: total=2.879, reg=0.245, containing_obj=
0.918, no_obj=0.753, cls=0.362
Epoch [19/50], Iter [150/157], Loss: total=2.842, reg=0.241, containing_obj
=0.900, no_obj=0.755, cls=0.361
```

```
Starting epoch 20 / 50
Learning Rate for this epoch: 0.001
Epoch [20/50], Iter [50/157], Loss: total=2.865, reg=0.244, containing_obj=
0.890, no_obj=0.811, cls=0.350
Epoch [20/50], Iter [100/157], Loss: total=2.816, reg=0.241, containing_obj
=0.876, no_obj=0.779, cls=0.345
Epoch [20/50], Iter [150/157], Loss: total=2.803, reg=0.240, containing_obj
=0.873, no_obj=0.774, cls=0.342
---Evaluate model on test samples---
```

```
100%|██████████| 1255/1255 [00:44<00:00, 28.36it/s]
```



```
---class aeroplane ap 0.5178394111955906---
---class bicycle ap 0.5920482660698859---
---class bird ap 0.5201035804962366---
---class boat ap 0.23536275470413345---
---class bottle ap 0.20375844874492555---
---class bus ap 0.5361999417416985---
---class car ap 0.567787259824586---
---class cat ap 0.6341958822919206---
---class chair ap 0.29711824675561227---
---class cow ap 0.35192256418430234---
---class diningtable ap 0.30101536335799317---
---class dog ap 0.5490534122664519---
---class horse ap 0.5727695541299401---
---class motorbike ap 0.5154927620846118---
---class person ap 0.4224381653945489---
---class pottedplant ap 0.1393448673075399---
---class sheep ap 0.19784476321785938---
---class sofa ap 0.39209188130178696---
---class train ap 0.5548672829813818---
---class tvmonitor ap 0.5139563804401719---
---map 0.4307605394245589---
19 [0.5178394111955906, 0.5920482660698859, 0.5201035804962366, 0.235362754
70413345, 0.20375844874492555, 0.5361999417416985, 0.567787259824586, 0.634
1958822919206, 0.29711824675561227, 0.35192256418430234, 0.3010153633579931
7, 0.5490534122664519, 0.5727695541299401, 0.5154927620846118, 0.4224381653
945489, 0.1393448673075399, 0.19784476321785938, 0.39209188130178696, 0.554
8672829813818, 0.5139563804401719]
```

Starting epoch 21 / 50

Learning Rate for this epoch: 0.001

Epoch [21/50], Iter [50/157], Loss: total=2.745, reg=0.241, containing_obj=0.854, no_obj=0.741, cls=0.315

Epoch [21/50], Iter [100/157], Loss: total=2.716, reg=0.236, containing_obj=0.854, no_obj=0.744, cls=0.313

Epoch [21/50], Iter [150/157], Loss: total=2.735, reg=0.236, containing_obj=0.852, no_obj=0.754, cls=0.328

Updating best val loss: 3.93851

Starting epoch 22 / 50

Learning Rate for this epoch: 0.001

Epoch [22/50], Iter [50/157], Loss: total=2.718, reg=0.239, containing_obj=0.835, no_obj=0.740, cls=0.321

Epoch [22/50], Iter [100/157], Loss: total=2.736, reg=0.238, containing_obj=0.852, no_obj=0.754, cls=0.317

Epoch [22/50], Iter [150/157], Loss: total=2.708, reg=0.234, containing_obj=0.846, no_obj=0.750, cls=0.318

Updating best val loss: 3.90568

Starting epoch 23 / 50

Learning Rate for this epoch: 0.001

Epoch [23/50], Iter [50/157], Loss: total=2.492, reg=0.216, containing_obj=0.782, no_obj=0.726, cls=0.266

Epoch [23/50], Iter [100/157], Loss: total=2.551, reg=0.218, containing_obj=0.809, no_obj=0.733, cls=0.284

Epoch [23/50], Iter [150/157], Loss: total=2.609, reg=0.223, containing_obj=0.825, no_obj=0.738, cls=0.300

Starting epoch 24 / 50

Learning Rate for this epoch: 0.001

Epoch [24/50], Iter [50/157], Loss: total=2.617, reg=0.221, containing_obj=

```
0.829, no_obj=0.765, cls=0.298
Epoch [24/50], Iter [100/157], Loss: total=2.630, reg=0.227, containing_obj=
0.822, no_obj=0.754, cls=0.293
Epoch [24/50], Iter [150/157], Loss: total=2.594, reg=0.225, containing_obj
=0.813, no_obj=0.741, cls=0.283
```

```
Starting epoch 25 / 50
Learning Rate for this epoch: 0.001
Epoch [25/50], Iter [50/157], Loss: total=2.545, reg=0.223, containing_obj=
0.790, no_obj=0.726, cls=0.278
Epoch [25/50], Iter [100/157], Loss: total=2.493, reg=0.216, containing_obj
=0.776, no_obj=0.732, cls=0.271
Epoch [25/50], Iter [150/157], Loss: total=2.505, reg=0.218, containing_obj
=0.780, no_obj=0.725, cls=0.275
---Evaluate model on test samples---
```

```
100%|██████████| 1255/1255 [00:44<00:00, 28.06it/s]
```

```
---class aeroplane ap 0.6618834802952756---
---class bicycle ap 0.5474977979359678---
---class bird ap 0.48342630900207656---
---class boat ap 0.35753590984435524---
---class bottle ap 0.2006726119240646---
---class bus ap 0.43443409407850947---
---class car ap 0.6481704403259481---
---class cat ap 0.6656904993642817---
---class chair ap 0.30660411942819715---
---class cow ap 0.4623193901826518---
---class diningtable ap 0.27730169709091534---
---class dog ap 0.6065831582261425---
---class horse ap 0.5826933866134423---
---class motorbike ap 0.6082032515809171---
---class person ap 0.47263671045605116---
---class pottedplant ap 0.19817411343375152---
---class sheep ap 0.35481881569183216---
---class sofa ap 0.4716556041879993---
---class train ap 0.6732552865099146---
---class tvmonitor ap 0.513116122402945---
---map 0.4763336399287619---
24 [0.6618834802952756, 0.5474977979359678, 0.48342630900207656, 0.35753590
984435524, 0.2006726119240646, 0.43443409407850947, 0.6481704403259481, 0.6
656904993642817, 0.30660411942819715, 0.4623193901826518, 0.277301697090915
34, 0.6065831582261425, 0.5826933866134423, 0.6082032515809171, 0.472636710
45605116, 0.19817411343375152, 0.35481881569183216, 0.4716556041879993, 0.6
732552865099146, 0.513116122402945]
```

Starting epoch 26 / 50

Learning Rate for this epoch: 0.001

Epoch [26/50], Iter [50/157], Loss: total=2.432, reg=0.212, containing_obj=0.746, no_obj=0.746, cls=0.253

Epoch [26/50], Iter [100/157], Loss: total=2.478, reg=0.214, containing_obj=0.772, no_obj=0.736, cls=0.268

Epoch [26/50], Iter [150/157], Loss: total=2.520, reg=0.218, containing_obj=0.785, no_obj=0.736, cls=0.277

Starting epoch 27 / 50

Learning Rate for this epoch: 0.001

Epoch [27/50], Iter [50/157], Loss: total=2.478, reg=0.212, containing_obj=0.795, no_obj=0.728, cls=0.257

Epoch [27/50], Iter [100/157], Loss: total=2.441, reg=0.208, containing_obj=0.778, no_obj=0.724, cls=0.262

Epoch [27/50], Iter [150/157], Loss: total=2.451, reg=0.211, containing_obj=0.779, no_obj=0.722, cls=0.257

Starting epoch 28 / 50

Learning Rate for this epoch: 0.001

Epoch [28/50], Iter [50/157], Loss: total=2.378, reg=0.206, containing_obj=0.758, no_obj=0.693, cls=0.243

Epoch [28/50], Iter [100/157], Loss: total=2.369, reg=0.205, containing_obj=0.759, no_obj=0.704, cls=0.236

Epoch [28/50], Iter [150/157], Loss: total=2.366, reg=0.203, containing_obj=0.753, no_obj=0.707, cls=0.246

Starting epoch 29 / 50

Learning Rate for this epoch: 0.001

Epoch [29/50], Iter [50/157], Loss: total=2.350, reg=0.201, containing_obj=0.726, no_obj=0.720, cls=0.259

Epoch [29/50], Iter [100/157], Loss: total=2.323, reg=0.201, containing_obj

```
=0.723, no_obj=0.709, cls=0.241  
Epoch [29/50], Iter [150/157], Loss: total=2.359, reg=0.205, containing_obj  
=0.737, no_obj=0.715, cls=0.240
```

```
Starting epoch 30 / 50  
Learning Rate for this epoch: 0.001  
Epoch [30/50], Iter [50/157], Loss: total=2.300, reg=0.193, containing_obj=  
0.726, no_obj=0.740, cls=0.237  
Epoch [30/50], Iter [100/157], Loss: total=2.347, reg=0.201, containing_obj  
=0.738, no_obj=0.731, cls=0.238  
Epoch [30/50], Iter [150/157], Loss: total=2.352, reg=0.201, containing_obj  
=0.739, no_obj=0.740, cls=0.236  
---Evaluate model on test samples---
```

```
100%|██████████| 1255/1255 [00:44<00:00, 28.15it/s]
```

```
---class aeroplane ap 0.5690340210774798---
---class bicycle ap 0.5190527848003978---
---class bird ap 0.5672158450048697---
---class boat ap 0.3222040025624234---
---class bottle ap 0.24144738633958504---
---class bus ap 0.5241597408320332---
---class car ap 0.6976395280141234---
---class cat ap 0.7012469064932182---
---class chair ap 0.32479549979349936---
---class cow ap 0.44406714178766327---
---class diningtable ap 0.2797851151439072---
---class dog ap 0.6677157348219405---
---class horse ap 0.6273468301531409---
---class motorbike ap 0.44962175727908---
---class person ap 0.4533645350526373---
---class pottedplant ap 0.15738595286061463---
---class sheep ap 0.281355942423871---
---class sofa ap 0.357736699166503---
---class train ap 0.6703515581943431---
---class tvmonitor ap 0.5684723384404881---
---map 0.47119996601209085---
29 [0.5690340210774798, 0.5190527848003978, 0.5672158450048697, 0.322204002
5624234, 0.24144738633958504, 0.5241597408320332, 0.6976395280141234, 0.701
2469064932182, 0.32479549979349936, 0.44406714178766327, 0.279785115143907
2, 0.6677157348219405, 0.6273468301531409, 0.44962175727908, 0.453364535052
6373, 0.15738595286061463, 0.281355942423871, 0.357736699166503, 0.67035155
81943431, 0.5684723384404881]
Updating best val loss: 3.89517
```

```
Starting epoch 31 / 50
Learning Rate for this epoch: 0.0001
Epoch [31/50], Iter [50/157], Loss: total=2.259, reg=0.197, containing_obj=
0.719, no_obj=0.680, cls=0.213
Epoch [31/50], Iter [100/157], Loss: total=2.231, reg=0.193, containing_obj
=0.719, no_obj=0.683, cls=0.205
Epoch [31/50], Iter [150/157], Loss: total=2.206, reg=0.190, containing_obj
=0.704, no_obj=0.685, cls=0.208
Updating best val loss: 3.78885
```

```
Starting epoch 32 / 50
Learning Rate for this epoch: 0.0001
Epoch [32/50], Iter [50/157], Loss: total=2.152, reg=0.183, containing_obj=
0.694, no_obj=0.675, cls=0.204
Epoch [32/50], Iter [100/157], Loss: total=2.099, reg=0.180, containing_obj
=0.671, no_obj=0.676, cls=0.191
Epoch [32/50], Iter [150/157], Loss: total=2.115, reg=0.183, containing_obj
=0.676, no_obj=0.669, cls=0.192
```

```
Starting epoch 33 / 50
Learning Rate for this epoch: 0.0001
Epoch [33/50], Iter [50/157], Loss: total=2.168, reg=0.189, containing_obj=
0.690, no_obj=0.688, cls=0.187
Epoch [33/50], Iter [100/157], Loss: total=2.140, reg=0.186, containing_obj
=0.677, no_obj=0.687, cls=0.190
Epoch [33/50], Iter [150/157], Loss: total=2.097, reg=0.182, containing_obj
=0.665, no_obj=0.680, cls=0.182
```

```
Starting epoch 34 / 50
Learning Rate for this epoch: 0.0001
Epoch [34/50], Iter [50/157], Loss: total=2.017, reg=0.171, containing_obj=
```

```
0.644, no_obj=0.681, cls=0.179
Epoch [34/50], Iter [100/157], Loss: total=2.064, reg=0.175, containing_obj=
0.665, no_obj=0.668, cls=0.190
Epoch [34/50], Iter [150/157], Loss: total=2.081, reg=0.178, containing_obj
=0.666, no_obj=0.673, cls=0.188
```

```
Starting epoch 35 / 50
Learning Rate for this epoch: 0.0001
Epoch [35/50], Iter [50/157], Loss: total=2.058, reg=0.176, containing_obj=
0.665, no_obj=0.683, cls=0.170
Epoch [35/50], Iter [100/157], Loss: total=2.044, reg=0.176, containing_obj
=0.645, no_obj=0.685, cls=0.175
Epoch [35/50], Iter [150/157], Loss: total=2.018, reg=0.173, containing_obj
=0.640, no_obj=0.680, cls=0.174
---Evaluate model on test samples---
```

```
100%|██████████| 1255/1255 [00:44<00:00, 28.39it/s]
```

```
---class aeroplane ap 0.5401731935960894---
---class bicycle ap 0.5954041400939147---
---class bird ap 0.5333325485578122---
---class boat ap 0.3937482947649102---
---class bottle ap 0.23302052145410085---
---class bus ap 0.5003197129834902---
---class car ap 0.706759231526782---
---class cat ap 0.7521723091486834---
---class chair ap 0.3467813809019478---
---class cow ap 0.4601272257689342---
---class diningtable ap 0.33357351593055967---
---class dog ap 0.6557478081154963---
---class horse ap 0.6800002750675859---
---class motorbike ap 0.5222948102817202---
---class person ap 0.4991196878808704---
---class pottedplant ap 0.19954851133647403---
---class sheep ap 0.3334870259996898---
---class sofa ap 0.44599772703153207---
---class train ap 0.781119370750939---
---class tvmonitor ap 0.5851359284523406---
---map 0.5048931609821936---
34 [0.5401731935960894, 0.5954041400939147, 0.5333325485578122, 0.393748294
7649102, 0.23302052145410085, 0.5003197129834902, 0.706759231526782, 0.7521
723091486834, 0.3467813809019478, 0.4601272257689342, 0.33357351593055967,
0.6557478081154963, 0.6800002750675859, 0.5222948102817202, 0.4991196878808
704, 0.19954851133647403, 0.3334870259996898, 0.44599772703153207, 0.781119
370750939, 0.5851359284523406]
Updating best val loss: 3.76522
```

Starting epoch 36 / 50

Learning Rate for this epoch: 0.0001

Epoch [36/50], Iter [50/157], Loss: total=1.976, reg=0.167, containing_obj=0.639, no_obj=0.661, cls=0.172

Epoch [36/50], Iter [100/157], Loss: total=1.995, reg=0.171, containing_obj=0.644, no_obj=0.651, cls=0.172

Epoch [36/50], Iter [150/157], Loss: total=2.022, reg=0.173, containing_obj=0.658, no_obj=0.645, cls=0.175

Starting epoch 37 / 50

Learning Rate for this epoch: 0.0001

Epoch [37/50], Iter [50/157], Loss: total=2.033, reg=0.176, containing_obj=0.646, no_obj=0.693, cls=0.160

Epoch [37/50], Iter [100/157], Loss: total=2.028, reg=0.179, containing_obj=0.630, no_obj=0.668, cls=0.168

Epoch [37/50], Iter [150/157], Loss: total=2.018, reg=0.177, containing_obj=0.629, no_obj=0.666, cls=0.174

Starting epoch 38 / 50

Learning Rate for this epoch: 0.0001

Epoch [38/50], Iter [50/157], Loss: total=2.041, reg=0.178, containing_obj=0.644, no_obj=0.675, cls=0.171

Epoch [38/50], Iter [100/157], Loss: total=2.040, reg=0.179, containing_obj=0.630, no_obj=0.672, cls=0.178

Epoch [38/50], Iter [150/157], Loss: total=2.006, reg=0.176, containing_obj=0.618, no_obj=0.677, cls=0.168

Starting epoch 39 / 50

Learning Rate for this epoch: 0.0001

Epoch [39/50], Iter [50/157], Loss: total=2.013, reg=0.172, containing_obj=0.659, no_obj=0.660, cls=0.164

Epoch [39/50], Iter [100/157], Loss: total=2.026, reg=0.173, containing_obj=0.648, no_obj=0.677, cls=0.173
Epoch [39/50], Iter [150/157], Loss: total=1.993, reg=0.170, containing_obj=0.635, no_obj=0.661, cls=0.177

Starting epoch 40 / 50

Learning Rate for this epoch: 0.0001

Epoch [40/50], Iter [50/157], Loss: total=1.936, reg=0.162, containing_obj=0.629, no_obj=0.647, cls=0.173

Epoch [40/50], Iter [100/157], Loss: total=1.924, reg=0.164, containing_obj=0.619, no_obj=0.653, cls=0.160

Epoch [40/50], Iter [150/157], Loss: total=1.955, reg=0.167, containing_obj=0.624, no_obj=0.660, cls=0.165

---Evaluate model on test samples---

100%|██████████| 1255/1255 [00:46<00:00, 27.16it/s]


```
---class aeroplane ap 0.6445665071585093---
---class bicycle ap 0.5807727318309939---
---class bird ap 0.5294340875989305---
---class boat ap 0.38203859577435095---
---class bottle ap 0.20810453609862578---
---class bus ap 0.4679455698474837---
---class car ap 0.6921615741842123---
---class cat ap 0.806354583132037---
---class chair ap 0.37258656515824684---
---class cow ap 0.47365256927985827---
---class diningtable ap 0.3122347649267324---
---class dog ap 0.6736621107023588---
---class horse ap 0.6730531670241466---
---class motorbike ap 0.4971324649895331---
---class person ap 0.508266678823903---
---class pottedplant ap 0.21517028871543967---
---class sheep ap 0.37828634298869174---
---class sofa ap 0.46522404849192556---
---class train ap 0.7540345444515948---
---class tvmonitor ap 0.5843622117553461---
---map 0.5109521971466461---
39 [0.6445665071585093, 0.5807727318309939, 0.5294340875989305, 0.382038595
77435095, 0.20810453609862578, 0.4679455698474837, 0.6921615741842123, 0.80
6354583132037, 0.37258656515824684, 0.47365256927985827, 0.312234764926732
4, 0.6736621107023588, 0.6730531670241466, 0.4971324649895331, 0.5082666788
23903, 0.21517028871543967, 0.37828634298869174, 0.46522404849192556, 0.754
0345444515948, 0.5843622117553461]
```

Starting epoch 41 / 50

Learning Rate for this epoch: 1e-05

Epoch [41/50], Iter [50/157], Loss: total=1.989, reg=0.176, containing_obj=0.613, no_obj=0.671, cls=0.163

Epoch [41/50], Iter [100/157], Loss: total=2.002, reg=0.176, containing_obj=0.625, no_obj=0.667, cls=0.165

Epoch [41/50], Iter [150/157], Loss: total=1.993, reg=0.174, containing_obj=0.624, no_obj=0.668, cls=0.163

Starting epoch 42 / 50

Learning Rate for this epoch: 1e-05

Epoch [42/50], Iter [50/157], Loss: total=1.926, reg=0.167, containing_obj=0.594, no_obj=0.658, cls=0.168

Epoch [42/50], Iter [100/157], Loss: total=1.944, reg=0.169, containing_obj=0.598, no_obj=0.654, cls=0.176

Epoch [42/50], Iter [150/157], Loss: total=1.977, reg=0.173, containing_obj=0.613, no_obj=0.651, cls=0.172

Updating best val loss: 3.76134

Starting epoch 43 / 50

Learning Rate for this epoch: 1e-05

Epoch [43/50], Iter [50/157], Loss: total=1.978, reg=0.174, containing_obj=0.610, no_obj=0.668, cls=0.165

Epoch [43/50], Iter [100/157], Loss: total=1.967, reg=0.173, containing_obj=0.606, no_obj=0.666, cls=0.163

Epoch [43/50], Iter [150/157], Loss: total=1.987, reg=0.173, containing_obj=0.625, no_obj=0.662, cls=0.166

Updating best val loss: 3.75632

Starting epoch 44 / 50

Learning Rate for this epoch: 1e-05

Epoch [44/50], Iter [50/157], Loss: total=1.873, reg=0.162, containing_obj=

```
0.592, no_obj=0.648, cls=0.149
Epoch [44/50], Iter [100/157], Loss: total=1.942, reg=0.164, containing_obj=
0.620, no_obj=0.654, cls=0.174
Epoch [44/50], Iter [150/157], Loss: total=1.943, reg=0.164, containing_obj
=0.623, no_obj=0.660, cls=0.170
```

```
Starting epoch 45 / 50
Learning Rate for this epoch: 1e-05
Epoch [45/50], Iter [50/157], Loss: total=2.012, reg=0.173, containing_obj=
0.650, no_obj=0.650, cls=0.173
Epoch [45/50], Iter [100/157], Loss: total=1.955, reg=0.169, containing_obj
=0.617, no_obj=0.657, cls=0.163
Epoch [45/50], Iter [150/157], Loss: total=1.941, reg=0.168, containing_obj
=0.610, no_obj=0.651, cls=0.167
---Evaluate model on test samples---
```

```
100%|██████████| 1255/1255 [00:44<00:00, 28.47it/s]
```

```
---class aeroplane ap 0.6291865457980178---
---class bicycle ap 0.5615867307808892---
---class bird ap 0.5259251846047481---
---class boat ap 0.3669550472014277---
---class bottle ap 0.20827355047135687---
---class bus ap 0.5496316303030001---
---class car ap 0.6770119205188168---
---class cat ap 0.8125604630438231---
---class chair ap 0.35799751740477953---
---class cow ap 0.5125449697672544---
---class diningtable ap 0.3178491265637037---
---class dog ap 0.6692945783281833---
---class horse ap 0.6616610195272735---
---class motorbike ap 0.4932249646445406---
---class person ap 0.5086574985624499---
---class pottedplant ap 0.23186776304809406---
---class sheep ap 0.406924284256662---
---class sofa ap 0.45818712911435067---
---class train ap 0.7282925398935043---
---class tvmonitor ap 0.5879711866561131---
---map 0.5132801825244494---
44 [0.6291865457980178, 0.5615867307808892, 0.5259251846047481, 0.366955047
2014277, 0.20827355047135687, 0.5496316303030001, 0.6770119205188168, 0.812
5604630438231, 0.35799751740477953, 0.5125449697672544, 0.3178491265637037,
0.6692945783281833, 0.6616610195272735, 0.4932249646445406, 0.5086574985624
499, 0.23186776304809406, 0.406924284256662, 0.45818712911435067, 0.7282925
398935043, 0.5879711866561131]
Updating best val loss: 3.75198
```

Starting epoch 46 / 50

Learning Rate for this epoch: 1e-05

Epoch [46/50], Iter [50/157], Loss: total=1.925, reg=0.168, containing_obj=0.594, no_obj=0.641, cls=0.173

Epoch [46/50], Iter [100/157], Loss: total=1.951, reg=0.168, containing_obj=0.607, no_obj=0.649, cls=0.177

Epoch [46/50], Iter [150/157], Loss: total=1.919, reg=0.166, containing_obj=0.599, no_obj=0.647, cls=0.168

Starting epoch 47 / 50

Learning Rate for this epoch: 1e-05

Epoch [47/50], Iter [50/157], Loss: total=1.997, reg=0.173, containing_obj=0.636, no_obj=0.644, cls=0.174

Epoch [47/50], Iter [100/157], Loss: total=1.965, reg=0.171, containing_obj=0.621, no_obj=0.649, cls=0.166

Epoch [47/50], Iter [150/157], Loss: total=1.952, reg=0.170, containing_obj=0.611, no_obj=0.651, cls=0.166

Starting epoch 48 / 50

Learning Rate for this epoch: 1e-05

Epoch [48/50], Iter [50/157], Loss: total=1.885, reg=0.162, containing_obj=0.583, no_obj=0.654, cls=0.164

Epoch [48/50], Iter [100/157], Loss: total=1.926, reg=0.165, containing_obj=0.610, no_obj=0.653, cls=0.165

Epoch [48/50], Iter [150/157], Loss: total=1.922, reg=0.165, containing_obj=0.608, no_obj=0.650, cls=0.162

Starting epoch 49 / 50

Learning Rate for this epoch: 1e-05

Epoch [49/50], Iter [50/157], Loss: total=2.017, reg=0.171, containing_obj=0.645, no_obj=0.668, cls=0.181

Epoch [49/50], Iter [100/157], Loss: total=2.002, reg=0.171, containing_obj=0.635, no_obj=0.666, cls=0.177

Epoch [49/50], Iter [150/157], Loss: total=1.963, reg=0.168, containing_obj=0.621, no_obj=0.663, cls=0.169

Starting epoch 50 / 50

Learning Rate for this epoch: 1e-05

Epoch [50/50], Iter [50/157], Loss: total=1.991, reg=0.169, containing_obj=0.639, no_obj=0.662, cls=0.173

Epoch [50/50], Iter [100/157], Loss: total=1.990, reg=0.172, containing_obj=0.625, no_obj=0.666, cls=0.171

Epoch [50/50], Iter [150/157], Loss: total=1.988, reg=0.173, containing_obj=0.622, no_obj=0.669, cls=0.167

---Evaluate model on test samples---

100%|██████████| 1255/1255 [00:45<00:00, 27.35it/s]

---class aeroplane ap 0.6022119783511177---

---class bicycle ap 0.5721199352176295---

---class bird ap 0.5272814077186743---

---class boat ap 0.37958515424242983---

---class bottle ap 0.1970295655483863---

---class bus ap 0.5388653729384907---

---class car ap 0.6964292289889491---

---class cat ap 0.7818175308873413---

---class chair ap 0.36967891274796927---

---class cow ap 0.4839731910388159---

---class diningtable ap 0.3087458017418613---

---class dog ap 0.6580636661769885---

---class horse ap 0.6655591059538684---

---class motorbike ap 0.4908202964998829---

---class person ap 0.5032313386719985---

---class pottedplant ap 0.21362909914368527---

---class sheep ap 0.4128880742075391---

---class sofa ap 0.49837665109447216---

---class train ap 0.7531128368339353---

---class tvmonitor ap 0.5869395004225786---

---map 0.5120179324213308---

49 [0.6022119783511177, 0.5721199352176295, 0.5272814077186743, 0.37958515424242983, 0.1970295655483863, 0.5388653729384907, 0.6964292289889491, 0.7818175308873413, 0.36967891274796927, 0.4839731910388159, 0.3087458017418613, 0.6580636661769885, 0.6655591059538684, 0.4908202964998829, 0.5032313386719985, 0.21362909914368527, 0.4128880742075391, 0.49837665109447216, 0.7531128368339353, 0.5869395004225786]

View example predictions

```
In [ ]: net.eval()

# select random image from val set
image_name = random.choice(val_dataset.fnames)
image = cv2.imread(os.path.join(file_root_val, image_name))
image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

print('predicting...')
result = predict_image(net, image_name, root_img_directory=file_root_val)
for left_up, right_bottom, class_name, _, prob in result:
    color = COLORS[VOC_CLASSES.index(class_name)]
    cv2.rectangle(image, left_up, right_bottom, color, 2)
    label = class_name + str(round(prob, 2))
    text_size, baseline = cv2.getTextSize(label, cv2.FONT_HERSHEY_SIMPLEX, 0.8, 1)
    p1 = (left_up[0], left_up[1] - text_size[1])
    cv2.rectangle(image, (p1[0] - 2 // 2, p1[1] - 2 - baseline), (p1[0] + text_size[0] + 2 // 2, p1[1] - 2 - baseline), color, 1)
```

```

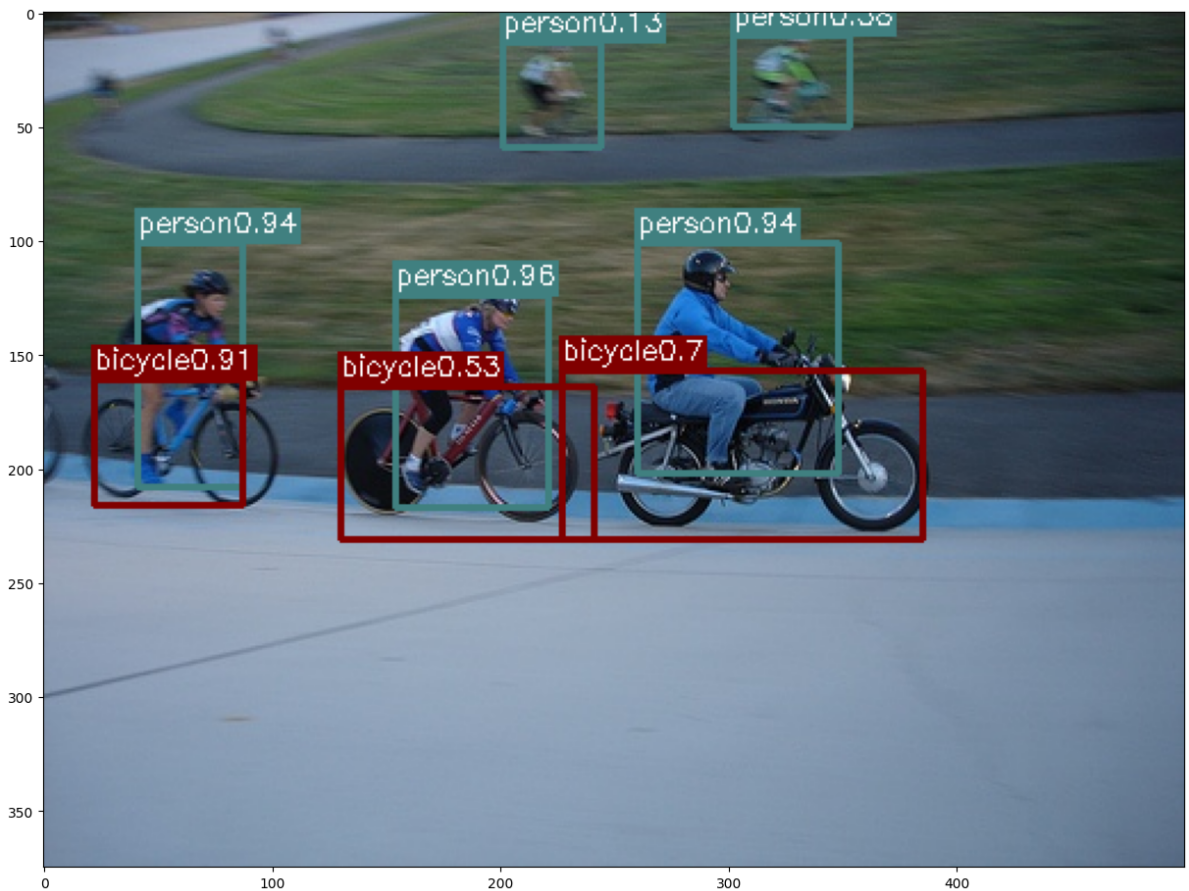
        color, -1)
    cv2.putText(image, label, (p1[0], p1[1] + baseline), cv2.FONT_HERSHEY_S:

plt.figure(figsize = (15,15))
plt.imshow(image)

```

predicting...

Out[]: <matplotlib.image.AxesImage at 0x7d09a7269f60>



Kaggle submission (85%)

Predict Result

Predict the results based on testing set. Upload to [Kaggle](#).

How to upload

1. Click the folder icon in the left hand side of Colab.
2. Right click "result.csv". Select "Download"
3. To kaggle. Click "Submit Predictions"
4. Upload the result.csv
5. System will automaticlaly calculate the accuracy of 50% dataset and publish this result to leaderboard.

預測 `test` 並將結果上傳至Kaggle。 [連結](#)

執行完畢此區的程式碼後，會將 `test` 預測完的結果存下來。

上傳流程

1. 點選左側選單最下方的資料夾圖示
2. 右鍵「result.csv」
3. 點選「Download」
4. 至連結網頁點選「Submit Predictions」
5. 將剛剛下載的檔案上傳
6. 系統會計算並公布其中50%資料的正確率

```
In [ ]: from google.colab import drive
drive.mount('/content/drive', force_remount=True)
```

```
In [ ]: %cd /content/drive/MyDrive/DeepLearning/A5
/content/drive/MyDrive/DeepLearning/A5
```

```
In [ ]: root_test = 'data/VOCdevkit_2007_2/VOC2007test/JPEGImages/'
file_test = 'data/voc2007test.txt'
```

By using the test_evaluate function, you will obtain predictions for each image.

```
In [ ]: preds_submission = test_evaluate(net, test_dataset_file=file_test, img_root=
---Evaluate model on test samples---
100%|██████████| 4950/4950 [43:41<00:00, 1.89it/s]
```

The write_csv function will use preds_submission to write into a CSV file called 'result.csv'.

```
In [ ]: write_csv(preds_submission)
```

Report (15%)

In your report, please include:

- a. A brief discussion on your implementation.
- b. Report the best train and validation accuracy in all of your experiments and discuss any strategies or tricks you've employed.
- c. Report the results for extra credits and also provide a discussion, if any.

Extra Credit (15%)

- Pick a fun video like [this one](#), run your detector on it (a subset of frames would be OK), and produce a video showing your results.
- Try to replace the provided pre-trained network with a different one and train with the YOLO loss on top to attempt to get better accuracy.
- Or any other methods that you try to improve the performance.

Video

Result link: <https://youtu.be/PN9ECezV37E>

```
In [ ]: ckpt = torch.load('/content/drive/MyDrive/DeepLearning/A5/checkpoints/best_ckpt.pth')
net.load_state_dict(ckpt)
```

```
In [ ]: video = cv2.VideoCapture("/content/Taylor Swift - We Are Never Ever Getting Back Together")
output_dir = "VideoData"

# Create the output directory if it doesn't exist
os.makedirs(output_dir, exist_ok=True)
frame_list = os.path.join(output_dir, 'frame_images_list.txt')
currentframe = 0

with open(frame_list, 'w') as frame_txt_f:
    while True:
        ret, frame = video.read()
        if ret:
            name = os.path.join(output_dir, 'frame' + str(currentframe) + '.jpg')
            # print('Creating...' + name)
            cv2.imwrite(name, frame)
            frame_txt_f.write(name + '\n')
            currentframe += 1
        else:
            break

video.release()
cv2.destroyAllWindows()
```

```
In [ ]: images_dir = "/content/drive/MyDrive/DeepLearning/A5/VideoData"
output_dir = "/content/drive/MyDrive/DeepLearning/A5/ResultData"
output_video_path = "/content/drive/MyDrive/DeepLearning/A5/video_result.mp4"
os.makedirs(output_dir, exist_ok=True)
```

```
In [ ]: image_path = os.path.join('VideoData', 'frame0.jpg')
image = cv2.imread(image_path)
image.shape
```

```
In [ ]: image_files = [f for f in os.listdir(images_dir)][1:]
fourcc = cv2.VideoWriter_fourcc(*'mp4v')
video_writer = cv2.VideoWriter(output_video_path, fourcc, 20.0, (1280, 720))
```

```
In [ ]: from tqdm import tqdm
from matplotlib import pyplot as plt
```

```
In [ ]: net.eval()
for image_file in tqdm(image_files, desc="Processing images"):
    image_path = os.path.join(images_dir, image_file)

    image = cv2.imread(image_path)

    result = predict_image(net, image_file, root_img_directory=f"{images_dir}")
    if result:
        for left_up, right_bottom, class_name, _, prob in result:
            color = COLORS[VOC_CLASSES.index(class_name)]
            cv2.rectangle(image, left_up, right_bottom, color, 2)
            label = class_name + str(round(prob, 2))
```

```
text_size, baseline = cv2.getTextSize(label, cv2.FONT_HERSHEY_S  
p1 = (left_up[0], left_up[1] - text_size[1])  
cv2.rectangle(image, (p1[0] - 2 // 2, p1[1] - 2 - baseline), (p1[0] + 2 // 2, p1[1] - 2 - baseline), (p1[0] + 2 // 2, p1[1] - 2 - baseline), color, -1)  
cv2.putText(image, label, (p1[0], p1[1] + baseline), cv2.FONT_HERSHEY_S  
  
output_image_path = os.path.join(output_dir, f'pred_{image_file}.jpg')  
cv2.imwrite(output_image_path, image)  
  
video_writer.write(image)  
  
video_writer.release()
```